

1 introduction

In science computer simulation has become an important tool. Simulation are becoming more and more advanced which increase the amount of data that are being generated. This data gets stored on harddrive and loaded again when its time to analyze the data. By doing in-situ real-time analysis where the data gets analyzed immediately after being generated. By doing computer simulation this way it may be possible to save time and hardware resources.

2 Hardware

The program have been tested on a server with the following hardware.

Motherboard: Supermicro X11DPU-Z+

CPU: Intel(R) Xeon(R) Gold 6130 CPU @ 2.10GHz, 32 core

DRAM: Samsung RDIMM, 2666 MT/s.

NVDIMM: Micron Technology NV-DIMM , 2933 MT/s

Both CPU have twelve memory slots each. Each CPU have six channels. There are one DRAM and one NVDIMM sharing one channel.

3 Coding with NVDIMM

3.1 Creating pmempool

Before NVDIMM can be used, the user must create whats called a memory pool on the NVDIMM. The NVDIMM has several modes, in order to be able to create a memory pool the mode must be set to fsdax. On a server this must be done by the system administrator. To see what mode the NVDIMM is in can be done by the command `ndctl-list`. A program called `pmempool` must also be installed on the server, it is this program that will create the memory pool. The command used for creating the memory pool for this thesis is

```
1 pmempool create --layout my_layout --size=50G obj pool.obj
```

The layout is a string stored in the memory pool. When a program access a memory pool it need to send a string that match the string in the memory pool in order to use it. The user can specify the size of the memory pool, if size is not specified the `pmempool` will create a pool

with the lowest size allowed. There are three different types of memory pool to choose from, they are obj, log and blk. Which type of memory pool to use depends on which type of library is used in the program. In this thesis the libpmemobj library was used and that is why obj was used in the creating of memory pool. For log and blk are for the libraries libpmemlog and libpmemblk. The last part of the command line is the name and file address of the memory pool.

3.2 Different types of libraries

The libpmemobj library is the main library that most people will use they are programming with NVDIMM.

libpmemlog
libpmemblk

3.3 Opening the memory pool

4 Benchmarks

There are four different benchmarks.

4.1 STREAM DRAM

The stream DRAM benchmark measure the memory speed of the DRAM. The benchmark uses the STREAM[1] benchmark without any changes in order to measure how fast the memory speed is. The benchmark run the test 32 times and only on one socket, every times it restart one extra thread is added. The result is as one would expect, adding more threads in beginning will give a big increase in transfer speed. But at thread 5 there gains in transfer speed will start to diminish and at thread 11 there will be very little increase in transfer speed when adding more threads.

4.2 STREAM NVDIMM

The stream NVDIMM benchmark measure the memory speed of the NVDIMM. This benchmark is the same as the STREAM DRAM benchmark mention above. The different is that the memory type have been changed from DRAM to NVDIMM. The original code looks like this.



Figure 1: DRAM Stream

Listing 1: Description

```

1 #ifndef STREAM_TYPE
2 #define STREAM_TYPE double
3 #endif
4
5 static STREAM_TYPE a[STREAM_ARRAY_SIZE+OFFSET],
6                   b[STREAM_ARRAY_SIZE+OFFSET],
7                   c[STREAM_ARRAY_SIZE+OFFSET];

```

It have been changed into this. In addition the PMEMobjpool must be initiated in main method.

```


1 PMEMobjpool *pop;
2 POBJ_LAYOUT_BEGIN(array);
3 POBJ_LAYOUT_TOID(array, double);
4 POBJ_LAYOUT_END(array);
5 TOID(double) a;
6 TOID(double) b;
7 TOID(double) c;
8
9 int main()
10 {
11     const char path[] = "/mnt/pmem0-xfs/pool.obj";
12     pop = pmemobj_create(path, LAYOUT_NAME, 10737418240,

```

```

0666);
13  if (pop == NULL)
14      pop = pmemobj_open(path, LAYOUT_NAME);
15
16      if (pop == NULL) {
17          perror(path);
18          exit(1);
19      }
20  }

```



4.3 benchmark 3

Graphs and tables below show the speed of a certain amount of NVDIMM threads while the rest of the threads are from DRAM to DRAM. The test have been conducted by transfer data simultaneously from DRAM-DRAM and NVM-NVM. All the threads are transferring the values of one array to another, all the arrays have 100 million elements of type double. This transfer happens 5000 times and the graphs shows the average of the first 200 iterations. This is done to ensure that all the threads can't finish early and make the remaining threads faster. The sum graphs shows the sum bandwidth of DRAM and NVM. Average graphs shows the average bandwidth of DRAM and NVM.

4.3.1 The code

The benchmark have three different program, the code is mostly the same except for the part where NVDIMM threads from NVDIMM-NVDIMM, DRAM-NVDIMM or NVDIMM-DRAM.

The entire code is run in the main function except when it finds the current time, that is done in a different function. The code begins with creating a memory pool and reads the parameters from command line. The parameters are how many threads are using the NVDIMM, the total amount of threads used and how many time the test will repeat itself. The code will then enter parallel area where one thread will create a 2d array where all the threads will save the time it took to copy the array.

Every threads will then create their two arrays. The thread id will determine if both arrays will be DRAM array, or if one or two arrays will be NVDIMM array. Both arrays will be DRAM if the thread id is lower than the total amount of threads minus the number of NVDIMM threads. Thread ids equal or higher than that will either have one or both arrays stored in NVDIMM. All the arrays in all of the threads will be populated by random numbers. When a thread is done it will wait at a barrier until all the other threads are populating their threads.

Listing 2: Creation of DRAM and NVDIMM arrays

```
1  if(thread_id < totalThreads-nvmThreads){
2      //From DRAM to DRAM
3      drm_read_array =
4          ((double*)malloc(ARRAY_LENGTH*sizeof(double)));
5      drm_write_array =
6          ((double*)malloc(ARRAY_LENGTH*sizeof(double)));
7      #pragma omp critical
8      {
9          for(i=0;i<ARRAY_LENGTH;i++){
10             drm_read_array[i] =
11                 ((double)rand()/((double)(RAND_MAX)));
12             drm_write_array[i] =
13                 ((double)rand()/((double)(RAND_MAX)));
14         }
15     }
16 }
17 else if(thread_id >= totalThreads-nvmThreads){
18     //From NVDIMM to NVDIMM
19     POBJ_ALLOC(pop, &nvm_read_array, double, sizeof(double))
```

```

        * ARRAY_LENGTH, NULL, NULL);
16 POBJ_ALLOC(pop, &nvm_write_array, double, sizeof(double)
        * ARRAY_LENGTH, NULL, NULL);
17 #pragma omp critical
18 {
19     for(i=0;i<ARRAY_LENGTH;i++){
20         D_RW(nvm_read_array)[i] =
            ((double)rand()/(double)(RAND_MAX));
21         D_RW(nvm_write_array)[i] =
            ((double)rand()/(double)(RAND_MAX));
22     }
23 }
24 }

```

Threads with DRAM arrays and threads with one or two NVDIMM array will split into their own part of the code with an if-sentence. All the threads will run the as many times as specified in the parameters and save the time each test takes in the 2d array that was created in the beginning. When they are done they will free up the memory and leave the parallel area.

Listing 3: Threads running their test.

```

1  if(thread_id < totalThreads-nvmThreads){
2      //From DRAM to DRAM
3      for(i=0;i<total_tests;i++){
4          //Time start
5          test_time[thread_id][i] = mysecond();
6          for(j=0;j<ARRAY_LENGTH;j++){
7              drm_write_array[j] = drm_read_array[j];
8          }
9          //Time stop.
10         test_time[thread_id][i] = mysecond() -
            test_time[thread_id][i];
11     }
12 }
13 else if(thread_id >= totalThreads-nvmThreads){
14     //From NVDIMM to NVDIMM
15     for(i=0;i<total_tests;i++){
16         //Time start
17         test_time[thread_id][i] = mysecond2();
18         for(j=0;j<ARRAY_LENGTH;j++){
19             D_RW(nvm_write_array)[j] = D_RO(nvm_read_array)[j];

```

```

20     //Time stop.
21     test_time[thread_id][i] = mysecond2() -
        test_time[thread_id][i];
22 }
23 }

```

The code will then print the entire 2d array where the time measurements are stored to the terminal. Each line represent all the test done by one thread. In the beginning of each line the code will add either DRAM if both arrays are stored on DRAM. Or NVM if one or both arrays are stored on NVDIMM. When the program is done printing it will exit.

4.3.2 NVM-NVM

The tables below shows the result of the benchmark where one group of threads transfer from DRAM-DRAM and another group from NVDIMM-NVDIMM. The test result in the tables are the transfer speed in MB/s. The first table shows the combined transfer speed of all threads that are copying from DRAM-DRAM and the second table shows the combined speed of all the threads that are copying from NVDIMM-NVDIMM.

The first line in the first table in figure 2 shows the combined transfer speed of the DRAM-DRAM copying where there are one thread copying from NVDIMM-NVDIMM. That means there are 15 threads that are copying from DRAM-DRAM. The first line in the second table shows the transfer speed of that one thread. The columns shows the test numbers. The column with name the "1-20" shows the average transfer speed of the first 20 tests. The third table in figure 3 shows the combined transfer speed of all the 16 threads in the first two tables. There is also a graph where all three tables are being represented.

One of the hopes by using NVDIMM and DRAM simultaneously was that there would be an increase in the transfer speed. But by comparing copy on figure 1 with the sum on figure 4 one can see that there has been no increase in transfer speed. Both graphs shows a transfer speed on around 65000 MB/s.

	NVM-NVM					
	DRAM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	33979.84	64179.89	64788.61	64951.02	64650.49	64447.29
2	33153.33	60854.29	62114.14	61912.29	61761.77	61872.50
3	32879.35	57643.15	58614.63	58759.60	58625.89	58423.27
4	26839.12	54508.48	55674.22	55682.26	55415.60	55367.41
5	25687.57	51076.28	52241.04	52083.85	51964.28	51955.32
6	24209.41	47721.43	48861.96	48786.82	48552.39	48656.74
7	22425.14	43819.63	45050.64	44975.15	44882.40	44672.94
8	20117.60	40489.02	41752.05	41501.44	41524.94	41286.00
9	17713.80	36687.45	37360.21	37385.79	37156.23	37232.82
10	15212.97	32731.79	33432.15	33379.32	33220.57	33296.06
11	12941.82	28177.13	28767.32	28814.29	28732.10	28833.90
12	9894.49	23814.89	24401.20	24408.00	24416.73	24525.18
13	7341.90	17595.41	18030.92	18359.65	17857.03	18532.65
14	4758.41	12750.63	13284.93	13351.80	13790.00	13569.13
15	2394.31	6834.38	6992.16	7515.33	7445.93	7353.06
	NVM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	5114.00	3289.73	3253.33	3250.62	3250.58	3248.86
2	10235.69	6599.16	6520.83	6507.40	6520.58	6500.57
3	15165.52	10119.48	9983.90	9984.04	9966.19	9979.90
4	20114.02	13554.15	13438.86	13435.23	13423.07	13416.59
5	24936.64	17103.23	16991.10	16958.34	16961.03	16933.81
6	29623.07	20662.40	20464.90	20437.07	20468.27	20438.51
7	33972.44	22367.64	22014.43	21953.45	21980.77	22105.95
8	37897.59	27397.05	27264.10	27141.45	27219.06	27083.62
9	41598.76	30795.78	30818.58	30922.06	30851.27	30427.67
10	44808.35	34539.87	34494.76	34500.89	34350.06	34203.29
11	48545.49	37723.50	37394.55	37303.10	37247.85	36857.64
12	51215.90	41988.10	41845.44	41750.50	41731.61	41533.89
13	54109.52	44058.49	43833.39	43794.52	43761.88	42853.98
14	57066.38	48613.84	48218.46	48161.35	48318.21	47953.55
15	58853.11	51816.83	51824.03	51488.10	51789.92	51100.06

Figure 2: NVM-NVM 1-100 iteration, 16 threads total, 3rd version

	SUM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	39093.84	67469.62	68041.93	68201.64	67901.07	67696.15
2	43389.02	67453.45	68634.96	68419.69	68282.35	68373.06
3	48044.87	67762.63	68598.53	68743.64	68592.08	68403.17
4	46953.14	68062.63	69113.08	69117.49	68838.67	68784.00
5	50624.21	68179.51	69232.14	69042.18	68925.31	68889.13
6	53832.48	68383.83	69326.86	69223.88	69020.66	69095.26
7	56397.58	66187.27	67065.07	66928.60	66863.17	66778.89
8	58015.19	67886.07	69016.15	68642.89	68744.00	68369.63
9	59312.56	67483.22	68178.79	68307.84	68007.50	67660.49
10	60021.32	67271.67	67926.91	67880.21	67570.63	67499.35
11	61487.31	65900.63	66161.87	66117.39	65979.96	65691.54
12	61110.39	65802.98	66246.64	66158.50	66148.34	66059.07
13	61451.42	61653.90	61864.31	62154.17	61618.91	61386.63
14	61824.79	61364.48	61503.39	61513.15	62108.21	61522.68
15	61247.42	58651.21	58816.19	59003.42	59235.85	58453.12

Figure 3: NVM-NVM 1-100 iteration, 16 threads total, 3rd version



Figure 4: NVM-NVM graph 1-20, 3rd version

4.3.3 NVM-DRAM

This benchmark is similar to the previous benchmark. The only difference is that some threads will transfer data from NVDIMM-DRAM instead of NVDIMM-NVDIMM.

	NVM-DRAM					
	DRAM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	32437.33	51729.43	53490.40	54419.38	56626.69	57295.70
2	31785.84	58694.61	59123.37	59135.90	58803.95	59160.35
3	27089.46	55323.72	55359.79	55449.94	55124.84	55079.11
4	26523.71	51565.09	51498.67	51032.22	50944.45	51155.26
5	23943.39	47390.53	47067.10	47291.86	46960.13	47119.51
6	21332.19	43205.17	43003.44	43143.48	42653.83	42962.85
7	18079.52	38801.31	38958.62	38995.27	38681.43	39000.71
8	18219.63	34652.20	34829.29	34582.21	34809.29	34763.03
9	15716.34	30244.85	30604.89	30504.82	30347.98	30531.90
10	12959.49	25984.54	26132.58	26116.43	26282.85	25981.46
11	11550.12	21756.54	21772.80	21935.54	21740.99	21776.18
12	8936.14	17175.26	17608.93	17216.47	17843.56	17341.46
13	6236.36	12779.91	13194.20	13003.21	13181.25	13119.76
14	3603.24	8669.61	8644.75	8731.08	8916.27	8621.84
15	1793.63	4250.55	4490.40	4267.18	4502.44	4260.46
	NVM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	2170.83	3635.35	3940.80	4127.15	4313.01	4234.35
2	4446.39	7927.66	7856.40	7977.52	7974.72	7800.26
3	6524.15	12018.48	11973.65	11978.96	11956.69	11868.26
4	8968.17	16245.16	16145.53	16224.22	16169.83	15979.78
5	11227.93	20004.18	20569.08	20226.56	20499.64	20114.15
6	13267.48	24362.23	24984.69	24557.47	24613.37	24477.62
7	15471.30	28739.42	29026.57	28754.67	28980.22	28589.24
8	18094.36	32833.01	33428.60	33380.75	33069.79	33033.76
9	20805.37	37368.77	37715.52	37628.86	37511.17	37290.54
10	22740.89	41319.93	41966.16	42139.10	42023.48	41966.67
11	25533.28	45907.61	46625.38	46449.12	46306.21	46181.20
12	28124.41	50396.18	50722.28	51065.63	50556.47	50735.26
13	30225.94	54592.24	55465.44	55422.83	55294.43	55153.29
14	32659.06	58440.82	59868.63	59725.88	59542.50	59570.12
15	35484.88	63168.26	64013.14	64235.39	63940.74	64110.34

Figure 5: NVM-DRAM 1-100 iteration, 3rd version

	SUM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	34608.16	55364.78	57431.19	58546.53	60939.70	61530.05
2	36232.23	66622.28	66979.77	67113.41	66778.68	66960.61
3	33613.61	67342.20	67333.44	67428.90	67081.53	66947.36
4	35491.88	67810.25	67644.20	67256.43	67114.28	67135.04
5	35171.32	67394.71	67636.19	67518.42	67459.77	67233.66
6	34599.67	67567.40	67988.12	67700.95	67267.20	67440.47
7	33550.82	67540.73	67985.19	67749.94	67661.65	67589.95
8	36313.99	67485.21	68257.89	67962.96	67879.08	67796.79
9	36521.71	67613.62	68320.41	68133.69	67859.15	67822.43
10	35700.38	67304.47	68098.74	68255.52	68306.33	67948.13
11	37083.40	67664.15	68398.18	68384.66	68047.19	67957.39
12	37060.55	67571.44	68331.21	68282.10	68400.02	68076.72
13	36462.30	67372.15	68659.64	68426.04	68475.68	68273.05
14	36262.30	67110.42	68513.38	68456.96	68458.77	68191.96
15	37278.51	67418.81	68503.54	68502.57	68443.18	68370.81

Figure 6: NVM-DRAM 1-100 iteration, 3rd version




Figure 7: NVM-DRAM graph 1-20, 3rd version

4.3.4 DRAM-NVM

This benchmark is also similar to the other two benchmarks. This time some of the threads will transfer from DRAM-NVDIMM.

	DRAM-NVM					
	DRAM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	18762.63	62678.30	63604.73	63196.58	63136.35	63599.61
2	27382.09	58321.86	59823.75	59850.20	59792.36	59717.87
3	25055.31	55650.92	55734.25	56127.18	55630.90	55935.54
4	24531.82	51279.78	52005.76	52282.87	51782.02	52177.37
5	23933.44	46555.25	48391.69	47941.57	48093.38	47839.01
6	18038.64	43432.14	44116.64	44225.90	43685.08	44094.70
7	16733.09	38827.55	39930.92	39448.07	39906.00	39483.39
8	14948.52	35285.18	35922.06	35737.11	35492.87	35796.48
9	12648.41	31228.15	31784.53	31492.60	31380.92	31454.26
10	11685.27	26830.40	27658.00	27610.40	27050.44	27021.04
11	9125.40	22445.38	23500.71	22790.09	22663.17	22910.98
12	6988.07	18442.56	18641.75	18849.18	18429.07	18354.52
13	5241.80	13668.18	14290.47	14034.79	14232.71	13744.87
14	3489.73	9087.65	9672.94	9461.29	9476.61	9276.00
15	1706.71	4273.06	4665.14	4342.46	4592.38	4610.87
	NVM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	3375.57	4094.13	3968.66	4013.27	4024.96	3937.43
2	5969.30	8738.95	8141.12	7939.10	8120.81	7879.02
3	14655.50	12710.47	12200.25	12147.26	12139.38	12124.44
4	12183.05	16846.06	16447.48	16153.87	16528.34	16083.88
5	15414.28	21516.95	20524.82	20689.27	20457.96	20526.30
6	23350.53	25835.47	24888.70	24740.28	25046.00	24538.40
7	24810.32	30042.63	29430.07	29410.65	29124.31	29382.89
8	25246.16	34903.09	33344.27	33607.80	33385.67	33247.98
9	29498.49	38826.97	37916.57	38188.18	37813.69	37219.56
10	36784.96	43563.46	42352.29	42525.81	42036.90	42070.58
11	39799.65	47967.83	46875.74	47224.03	46257.64	46424.18
12	38783.93	52752.69	51908.85	51408.20	51339.13	50338.82
13	43481.73	56883.56	56624.74	55455.89	55879.32	56058.91
14	49369.10	61602.79	60907.89	60678.79	60065.98	60171.88
15	49133.33	67153.54	66972.79	63557.38	63874.03	63521.86

Figure 8: DRAM-NVM 1-100 iteration, 3rd version

	SUM					
Nvm-threads	1	1-20	21-40	41-60	61-80	81-100
1	22138.20	66772.43	67573.40	67209.85	67161.31	67537.04
2	33351.39	67060.81	67964.86	67789.29	67913.17	67596.89
3	39710.81	68361.39	67934.50	68274.44	67770.28	68059.98
4	36714.87	68125.85	68453.24	68436.74	68310.35	68261.24
5	39347.72	68072.20	68916.50	68630.84	68551.33	68365.32
6	41389.17	69267.61	69005.34	68966.18	68731.08	68633.10
7	41543.41	68870.19	69360.99	68858.72	69030.31	68866.28
8	40194.68	70188.28	69266.33	69344.90	68878.54	69044.46
9	42146.90	70055.12	69701.09	69680.78	69194.61	68673.83
10	48470.23	70393.86	70010.29	70136.21	69087.34	69091.62
11	48925.05	70413.21	70376.45	70014.13	68920.81	69335.17
12	45772.00	71195.25	70550.60	70257.38	69768.20	68693.34
13	48723.53	70551.74	70915.22	69490.69	70112.04	69803.78
14	52858.83	70690.44	70580.83	70140.07	69542.59	69447.88
15	50840.04	71426.60	71637.93	67899.84	68466.41	68132.72

Figure 9: DRAM-NVM 1-100 iteration, 3rd version




Figure 10: DRAM-NVM graph 1-20, 3rd version

5 Simulation

5.1 DRAM only

Threads	Total time	Iteration time	Calculation time
1	1217.624747	747.345183	470.279564
2	641.937113	399.266812	242.670301
3	442.936486	275.083333	167.853153
4	339.788006	211.632944	128.155062
5	276.30993	171.953573	104.356358
6	233.895282	146.63456	87.260722
7	204.532469	129.058858	75.473611
8	186.792502	118.626586	68.165916
9	174.884702	111.713336	63.171366
10	164.146195	107.15154	56.994655
11	156.103916	103.946944	52.156971
12	151.490592	102.12569	49.364902
13	147.958949	100.005358	47.953591
14	143.513119	98.887134	44.625985
15	140.027579	98.30932	41.718259
16	137.833572	98.451276	39.382296

Table 1: Simulation with only DRAM, n=16M

The generation of data happens inside a while-loop that will repeat the generation and analyzing of data 5000 times. The total time is measured by taking the time before and after the while-loop. The analyzing time is measured the same way by taking the time before and after the analyzing part of the program. Data generation time is measured by subtracting the analyze time from the total time. All the time measurements are made inside a pragma omp single so there are no reason to be worried if the mysecond-method is thread-safe. The mysecond-method have been copied from the original stream benchmark. All the arrays used in the code have been subjected to first touch before the time is measured.

Listing 4: while-loop

```
1 #pragma omp single
2 {
3     data_generation_time, = mysecond();
4 }
5 while( n<5000 ){
6     #pragma omp barrier
7     #pragma omp single
```

```

8   {
9       n++;
10      diff=0.0;
11      average = 0;
12      //completes the first part of the formula.
13      Wk_1_product = (omd + (d*Wk_1))*iN;
14  }
15
16  Data generation, see chapter 5.1.1
17
18  #pragma omp barrier
19  #pragma omp single
20  {
21      temp_x = xk_1;
22      xk_1 = x;
23      x = temp_x;
24      //starting time measurement of calculation.
25      temp_calc=mysecond();
26  }
27
28  Analyzing the data, see chapter 5.1.2
29
30  #pragma omp barrier
31  #pragma omp single
32  {
33      average *= iN;
34      analyze_time+=mysecond()-temp_calc;
35  }
36  }
37  #pragma omp single
38  {
39      data_generation_time, = mysecond() -
        data_generation_time;
40  }

```

Listing 5: while-loop

```

1  double mysecond() {
2      struct timeval tp;
3      struct timezone tzp;
4      int i;
5      i = gettimeofday(&tp,&tzp);
6      return ( (double) tp.tv_sec + (double) tp.tv_usec *

```

```

1.e-6 );
7 }

```

5.1.1 Data generation

The arrays `x` and `xk_1` are double arrays that have the length of 4'000'000 elements. `CRS_row_ptr` and `CRS_col_idx` are int array that have the length of 31'976'004 elements, `CRS_values` have the same length and a double array. The code run through the `x`-array one time and `xk_1` twice. It also runs through `CRS_row_ptr`, `CRS_col_idx` and `CRS_values` once. The formula for calculating memory traffic is $3 \cdot 4'000'000 + 2 \cdot 31'976'004$ this would amount to 607 MB per iteration of the while-loop.

Listing 6: Generation of data.

```

1 #pragma omp for reduction(max:diff)
2 for( i=0; i<nodes; i++){
3     x[i] = 0;
4     for( j=CRS_row_ptr[i]; j<CRS_row_ptr[i+1]; j++)
5         x[i] += CRS_values[j] * xk_1[CRS_col_idx[j]];
6     x[i] *= d;
7     x[i] += Wk_1_product;
8     //Comuting the difference between x^k and x^k-1
9     //and adds the biggest diff to diffX[thread_id]
10    if( x[i]-xk_1[i] > diff ){
11        diff = x[i] - xk_1[i];
12    }
13 }

```

Calculation					
		Factor	repeated	total	total MB
n	16,000,000	2.5	5000	200,000,000,000	1,600,000.00
edges	127,952,004	1.5	5000	959,640,030,000	7,677,120.24
	Benchmark	Predicted	Data generated		
Threads	DRAM speed	time	time measurement		
1	13265.2	699.36	747.35		
2	26431.7	350.98	399.27		
3	38557.2	240.61	275.08		
4	49447.8	187.61	211.63		
5	58565.4	158.41	171.95		
6	62627.3	148.13	146.63		
7	66290.8	139.95	129.06		
8	69130.7	134.20	118.63		
9	71409.6	129.91	111.71		
10	73271.5	126.61	107.15		
11	74191.2	125.04	103.95		
12	74211.4	125.01	102.13		
13	74659.4	124.26	100.01		
14	74715.4	124.17	98.89		
15	74486	124.55	98.31		
16	74328.7	124.81	98.45		

Table 2: Prediction of time taken for data generation, n=16M

5.1.2 Analyze

The analyze part run through the nodes array five times and add all the elements to the average variable. The average variable is divided by the number of elements in array, this is shown in the code in chapter 5.1. The memory traffic will amount of 160MB per while-loop iteration.

Listing 7: Analyzing the data.

```
1 //Analyse part
2 #pragma omp for reduction(+ : average)
3     for(i=0;i<nodes;i++){
4         average += xk_1[i];
5     }
6 #pragma omp for reduction(+ : average)
7     for(i=0;i<nodes;i++){
8         average += xk_1[i];
9     }
10 #pragma omp for reduction(+ : average)
11     for(i=0;i<nodes;i++){
12         average += xk_1[i];
13     }
14 #pragma omp for reduction(+ : average)
15     for(i=0;i<nodes;i++){
16         average += xk_1[i];
17     }
18 #pragma omp for reduction(+ : average)
19     for(i=0;i<nodes;i++){
20         average += xk_1[i];
21     }
```

Calculation For analyze					
		Factor	MB	repeated	total MB
n	16,000,000.00	5	640	5000	3,200,000
	Benchmark		Analyze		Analyze time
Threads	DRAM speed		Prediction		Measurement
1	6665.4		480.09		470.279564
2	13183.1		242.74		242.670301
3	18704.7		171.08		167.853153
4	24971.9		128.14		128.155062
5	31169.6		102.66		104.356358
6	37259.5		85.88		87.260722
7	42771.7		74.82		75.473611
8	47211.9		67.78		68.165916
9	50706.4		63.11		63.171366
10	56235.2		56.90		56.994655
11	61401		52.12		52.156971
12	65062		49.18		49.364902
13	66645		48.02		47.953591
14	71518.7		44.74		44.625985
15	76423.3		41.87		41.718259
16	81259.4		39.38		39.382296

Table 3: Prediction of time taken for analyzing the data, n=16M

5.1.3 Stream benchmark, sum

This benchmark is the STREAM benchmark with an added benchmark. The sum is found by adding all the elements into a single variable. The STREAM benchmark was changed by increasing several array from four to five and added the sum benchmark after the four other benchmark. The STREAM benchmark for NVDIMM is the same as the one described above, but the code has been changed so the benchmark will read and write to the NVDIMM.

	copy	scale	add	triad	Average
1	13072.2	14053.7	13810.8	13821.3	6652.3
2	24586.2	25448.2	26622	25639.9	12920.8
3	35144.7	36424.1	38554.9	38323.4	18936.7
4	45261.6	46928	49559	49473.3	25020.5
5	52061.5	52960.9	58313.3	58783.2	31049.9
6	55459.1	55960	62524.6	62896.5	37123.1
7	57928.9	58281.5	65642	66144.9	42689.6
8	59714.4	60302.5	68166.4	68608.6	47295.1
9	61802.5	62368.3	70859.2	71188.9	50991.5
10	63482.1	63966.8	72520.9	72851	56010.9
11	64636.5	65069.6	73669.4	73741.7	61519.2
12	65450.4	65659.8	74202.1	74163.3	64427.4
13	66145.8	66209.1	74432.5	74503.6	66549.8
14	66337.4	66343.3	74314.9	74386.3	71372.7
15	66439.8	66274.5	74195	74232.2	76140.7
16	66417.5	66115.8	74067.2	74154.5	80702.4

Table 4: New Stream benchmark, DRAM

NVM sum test	cpu 0-15				
threads	copy	scale	add	triad	Sum
1	6455.2	4492.9	6134.2	5451.6	5422.2
2	12606.9	8783.7	11978.3	10631.0	10423.9
3	18284.2	12732.8	17392.6	15414.4	15141.5
4	24033.7	16760.6	22869.3	20280.1	19841.8
5	29584.8	20683.6	28242.3	25052.1	24489.6
6	35355.3	24818.6	33865.8	30064.3	29347.0
7	37564.9	24631.4	29185.7	26727.6	34074.4
8	45400.6	31945.8	43520.9	38752.7	37750.0
9	48372.0	33983.1	46176.0	41174.1	40607.1
10	52953.8	37741.9	51064.9	45665.5	45129.2
11	56689.4	41441.1	55640.6	50013.6	49416.7
12	58537.3	43532.5	56954.3	51616.4	52280.9
13	59904.7	44258.9	59200.8	53467.5	53727.5
14	62678.7	47558.2	63158.1	57325.7	57790.7
15	50577.2	36717.5	49718.4	44182.5	45334.0
16	66047.5	52866.6	66134.0	60444.3	65552.6

Table 5: New Stream benchmark, NVM

5.1.4 Calculation only

Calculation					
		Factor	repeated	total	total MB
n	16,000,000	2.5	5000	200,000,000,000	1,600,000.00
edges	127,952,004	1.5	5000	959,640,030,000	7,677,120.24
	Benchmark	Predicted	Data generated		
Threads	DRAM speed	time	time measurement		
1	13265.2	699.36	727.54		
2	26431.7	350.98	383.48		
3	38557.2	240.61	268.12		
4	49447.8	187.61	208.36		
5	58565.4	158.41	168.77		
6	62627.3	148.13	146.18		
7	66290.8	139.95	126.40		
8	69130.7	134.20	116.22		
9	71409.6	129.91	110.54		
10	73271.5	126.61	105.72		
11	74191.2	125.04	102.53		
12	74211.4	125.01	101.56		
13	74659.4	124.26	99.68		
14	74715.4	124.17	98.92		
15	74486	124.55	99.10		
16	74328.7	124.81	98.76		

Table 6: Prediction of time taken for calculation of the data, with a code that only do calculation

5.2 NVDIMM Analyze only

1	912.95
2	417.38
3	277.17
4	217.33
5	174.67
6	188.40
7	233.93
8	234.64
9	220.51
10	207.10
11	188.31
12	188.65
13	174.08
14	168.63
15	161.06
16	162.10

Table 7: Mearsurement of analyzations only on NVDIMM.

5.3 NVM simulation

5.3.1 Locks

The program are divided into two parts, the calculation of data and the analyzing of the data that have been generated. The two parts synchronize by using two locks called lock_a and lock_b, lock_a will start in unlocked state and lock_b in locked state. When the two parts starts the analyzing part is put on hold by lock_b until the calculation part has generated the first set of data. Then the calculation part will lock lock_a, swap the pointers x and xk_1 and then unlock lock_b. The calculation part will then start the calculation of the next set of data, but wont swap pointers until analyzing part has transferred the content in xk_1 to NVDIMM and unlocked lock_a.

When the calculation part unlocks lock_b the analyzing will start transferring data from xk_1 to NVDIMM and unlock lock_a when it's done with the transfer. The analyzing part will then start analyzing the data on NVDIMM. When its done it will encounter lock_b and will wait there until calculation has a new set of data ready and has swapped the pointers.

Before and after the set lock in calculation and analyze part there is a time measurement that measure how long the threads have waited for the lock to be unlocked by the other part. All the individual times the threads have waited in calculation or analyze gets added to a variable called iteration_idle_time or transfer_idle_time that will be the total time the threads have waited.

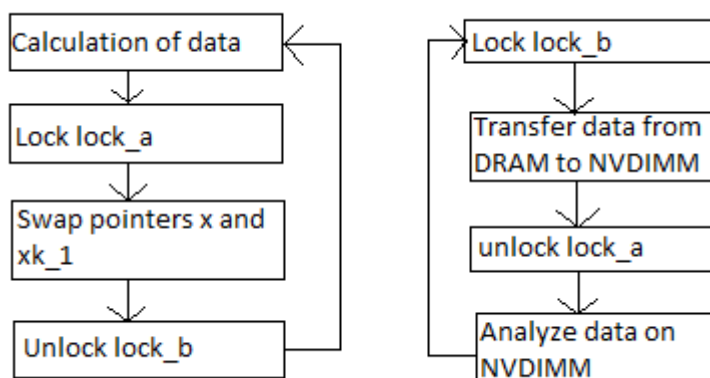


Figure 11: A simplified version of how lock works.

Listing 8: Calculation

```
1 while( n<5000 ){
2     #pragma omp barrier
3     /*
4         Calculation of Data
5     */
6     #pragma omp single
7     {
8         temp_time = mysecond();
9         omp_set_lock(&lock_a);
10        iteration_idle_time += mysecond() - temp_time;
11
12        temp_x = xk_1;
13        xk_1 = x;
14        x = temp_x;
15
16        omp_unset_lock(&lock_b);
17    }
18 } //end of while-loop
```

Listing 9: Analyze

```
1 while(1==1){
2     #pragma omp single
3     {
4         temp_time = mysecond();
5         omp_set_lock(&lock_b);
6         transfer_idle_time += mysecond() - temp_time;
7         temp_time = mysecond();
8         average=0.0;
9     }
10    /*
11        Transfer of array from DRAM to NVDIMM
12    */
13    #pragma omp single
14    {
15        DRAM_to_NVM_time += mysecond() - temp_time;
16        omp_unset_lock(&lock_a);
17        temp_time = mysecond();
18    }
19    /*
20        Analyzations of data
21    */
```

```

22     }
23     #pragma omp barrier
24     #pragma omp single
25     {
26         Analyse_time += mysecond() - temp_time;
27     }
28     //if sentence for exiting while-loop.
29     if(iteration_ongoing==0){
30         break;
31     }
32 }

```

Total	DataGen	Analyze	DataGen	DataGen	Transfer	DRAM-NVM	Analyze	Total
Threads	Threads	Threads	Time	Idle Time	Idle Time	Time	Time	Time
16	15	1	106.48	851.99	0.04	170.84	787.73	958.63
16	14	2	107.57	396.76	0.03	86.31	418.06	504.42
16	13	3	111.85	249.67	0.03	61.27	300.26	361.58
16	12	4	137.93	300.53	0.03	167.58	270.87	438.50
16	11	5	147.57	218.52	0.03	142.66	223.43	366.14
16	10	6	162.87	158.12	0.03	130.22	190.75	321.03
16	9	7	171.85	90.88	0.03	102.66	160.03	262.76
16	8	8	199.35	51.40	0.66	95.37	154.62	250.78
16	7	9	235.15	42.02	9.91	105.50	161.31	277.18
16	6	10	235.51	31.83	35.55	93.38	138.33	267.39
16	5	11	240.33	15.70	19.51	89.37	147.13	256.06
16	4	12	288.71	0.73	73.63	83.42	132.35	289.49
16	3	13	351.75	0.00	177.44	70.45	103.91	351.81
16	2	14	464.62	0.00	338.75	48.58	77.36	464.72
16	1	15	823.71	0.00	738.35	21.70	63.65	823.76

Table 8: Simulation with both NVDIMM and DRAM

getconf LEVEL1_DCACHE_LINESIZE have been used to find the cacheline size in bytes.

Predicted calculation	Load/store instructions completed	In MB	speed	seconds
15	161,347,231,158	1,290,777.85	74486.0	17.33
14	172,773,952,639	1,382,191.62	74715.4	18.50
13	185,957,951,084	1,487,663.61	74659.4	19.93
12	201,671,615,621	1,613,372.92	74211.4	21.74
11	219,943,929,111	1,759,551.43	74191.2	23.72
10	241,801,250,630	1,934,410.01	73271.5	26.40
9	268,439,799,619	2,147,518.40	71409.6	30.07
8	301,317,613,347	2,410,540.91	69130.7	34.87
7	343,673,207,234	2,749,385.66	66290.8	41.47
6	400,700,828,066	3,205,606.62	62627.3	51.19
5	480,131,664,993	3,841,053.32	58565.4	65.59
4	599,842,341,666	4,798,738.73	49447.8	97.05
3	799,776,584,260	6,398,212.67	38557.2	165.94
2	1,199,717,220,704	9,597,737.77	26431.7	363.11
1	2,399,341,092,964	19,194,728.74	13265.2	1447.00

Table 9: Predicted calculation based on papi.

threads	L3 cache misses	L3 cache misses	L3 load misses	L3 load misses
	<u>PAPI_L3_TCM</u> in MB	in MB	<u>PAPI_L3_LDM</u> in MB	in MB
15	9,702,008,094	620,928.52	25,760,001	1,648.64
14	10,414,505,930	666,528.38	24,990,133	1,599.37
13	11,267,940,351	721,148.18	18,020,986	1,153.34
12	12,187,448,519	779,996.71	18,274,879	1,169.59
11	13,274,582,641	849,573.29	18,637,339	1,192.79
10	14,593,925,871	934,011.26	20,253,335	1,296.21
9	16,197,929,846	1,036,667.51	22,522,674	1,441.45
8	18,226,402,006	1,166,489.73	24,733,748	1,582.96
7	20,895,036,560	1,337,282.34	26,738,318	1,711.25
6	24,379,213,623	1,560,269.67	30,061,653	1,923.95
5	29,243,830,364	1,871,605.14	36,954,105	2,365.06
4	36,569,741,426	2,340,463.45	44,680,152	2,859.53
3	48,621,709,482	3,111,789.41	56,447,194	3,612.62
2	72,963,408,589	4,669,658.15	80,746,396	5,167.77
1	146,650,902,016	9,385,657.73	149,083,280	9,541.33

Table 10: Tried to split load and store misses. Server do not support store misses.

	arrays	size	While-loop	sum
int	6	16,000,000	5000	480,000,000,000
int	2	127,952,004	5000	1,279,520,040,000
double	1	127,952,004	5000	639,760,020,000
			sum	2,399,280,060,000

Table 11: Manually counted the number of int and double transfers.

	PAPI_L3_TCM: 145437044007 cache misses.			
	PAPI_L3_TCW: 239627393 total cache writes.			
	Load/store	In MB	speed	seconds
	instructions completed			
1	145,676,671,400	9,323,306.97	13265.2	702.84

Table 12: Predicted calculationbased on papi with L3.

Calculation					
		Factor	repeated	total	total MB
n	16,000,000	2.5	5000	200,000,000,000.00	1,600,000.00
edges	127,952,004	1.5	5000	959,640,030,000.00	7,677,120.24
	Benchmark	Predicted	Data generated		
<u>Nvm-threads</u>	DRAM speed	time	time measurement		
1	74486.0	124.55	106.48		
2	74715.4	124.17	107.57		
3	74659.4	124.26	111.85		
4	74211.4	125.01	137.93		
5	74191.2	125.04	147.57		
6	73271.5	126.61	162.87		
7	71409.6	129.91	171.85		
8	69130.7	134.20	199.35		
9	66290.8	139.95	235.15		
10	62627.3	148.13	235.51		
11	58565.4	158.41	240.33		
12	49447.8	187.61	288.71		
13	38557.2	240.61	351.75		
14	26431.7	350.98	464.62		
15	13265.2	699.36	823.71		

Table 13: Time prediction, data generation

Analyze						
		Factor	repeated	total	total MB	
n	16,000,000	5	5000	400,000,000,000.00	3,200,000.00	
DRAM-NVM	16,000,000	1	5000	80,000,000,000.00	640,000.00	
	Benchmark	DRAM-NVM	DRAM-NVM	Benchmark	Analyze	Analyze time
<u>Nvm-threads</u>	<u>NVM speed</u>	Predicted	transfer time	<u>NVM sum</u>	Prediction	Measurement
1	3937.43	162.54	170.84	5449	587.26	787.73
2	7879.02	81.23	86.31	10216.5	313.22	418.06
3	12124.44	52.79	61.27	14921.8	214.45	300.26
4	16083.88	39.79	167.58	19562.3	163.58	270.87
5	20526.30	31.18	142.66	24142.1	132.55	223.43
6	24538.40	26.08	130.22	28837.2	110.97	190.75
7	29382.89	21.78	102.66	33421.1	95.75	160.03
8	33247.98	19.25	95.37	37228.5	85.96	154.62
9	37219.56	17.20	105.50	40080.3	79.84	161.31
10	42070.58	15.21	93.38	44540.9	71.84	138.33
11	46424.18	13.79	89.37	48703.0	65.70	147.13
12	50338.82	12.71	83.42	51520.0	62.11	132.35
13	56058.91	11.42	70.45	53418.7	59.90	103.91
14	60171.88	10.64	48.58	57326.6	55.82	77.36
15	63521.86	10.08	21.70	61542.9	52.00	63.65

Table 14: Time prediction, transfer and analyze

5.4 2D-array test

Listing 10: Kildekode

¹ https://github.com/SveinGunnar/Master_Thesis_2020/tree/master/ArrayCopyTest

4000*4000		
Threads	Time	Speed
1	138.47	4621.91
2	71.90	8901.87
3	49.98	12805.97
4	37.97	16854.35
5	30.83	20759.23
6	26.20	24426.96
7	23.37	27388.78
8	22.03	29050.67
9	21.13	30290.55
10	20.28	31558.77
11	19.77	32372.93
12	19.53	32769.09
13	19.95	32072.81
14	18.87	33919.44
15	18.83	33997.04
16	18.90	33864.18
10000*10000		
Threads	Time	Speed
1	856.44	4670.51
2	445.47	8979.21
3	306.34	13057.28
4	237.02	16876.45
5	195.01	20511.92
6	165.48	24172.77
7	153.00	26143.97
8	147.82	27059.68
9	146.34	27334.50
10	137.83	29021.36
11	134.70	29695.01
12	133.70	29917.48
13	131.01	30533.18
14	127.94	31263.72
15	128.40	31153.04
16	128.13	31217.93

Table 15: 2D-Array test

6 Large array benchmark

The size of data that must be analyzed keeps increasing year after year and the prize for DRAM are not getting cheaper. NVDIMM offer a lot of storage at a cheaper prize. This opens the opportunity to save money by offloading some of the data to the NVDIMM where the data will be analyzed the same way as the data on the DRAM. The downside to this strategy is that NVDIMM is slower than DRAM so the question is how much data can be offloaded to NVDIMM. If the user offload too much data to NVDIMM then the threads working on analyzing the data on DRAM will be idle while waiting for NVDIMM threads to complete.

Calculation

This program have an two dimensional array filled with data. The program start at element (1,1) of the array where it sum ups all of its eight neighbors and then takes the average. The result is stored in the same position in another two dimensional array. The program does this for every element between (1,1) and (m-2,n-2).

Formula

$$\frac{dram_data - nvdim_data}{dram_speed} = \frac{nvdim_data}{nvdim_speed}$$

$$nvdim_data = \frac{nvdim_speed * dram_data}{nvdim_speed + dram_speed}$$

Total_size	cores	dram_speed	cores	nvdim	nvdim_size	rows
16,000	10	48,656	6	20,438	4733	296
16,000	11	51,955	5	16,933	3933	246
16,000	12	55,367	4	13,416	3121	195
16,000	13	58,423	3	9,979	2334	146
16,000	14	61,872	2	6,500	1521	95
16,000	15	64,447	1	3,248	768	48

Table 16: Formula calculation

m	n	nvdim length	dram threads	nvdim threads	dram average	dram min	dram max	nvdim average	nvdim min	nvdim max	total average	total min	total max
1000	1000000	296	10	6	0.4040	0.3692	0.4309	0.9305	0.7452	0.7925	0.9305	0.7452	0.7925
1000	1000000	246	11	5	0.4078	0.3437	0.4417	0.9157	0.7431	0.7872	0.9158	0.7431	0.7872
1000	1000000	195	12	4	0.4120	0.3646	0.4522	0.8392	0.7250	0.8243	0.8392	0.7250	0.8243
1000	1000000	146	13	3	0.4381	0.3815	0.4868	0.8092	0.7243	0.7723	0.8092	0.7243	0.7724
1000	1000000	95	14	2	0.4359	0.4087	0.4752	0.7675	0.6902	0.7632	0.7675	0.6902	0.7632
1000	1000000	48	15	1	0.4459	0.3912	0.4722	0.6571	0.6014	0.6441	0.6571	0.6015	0.6441

Table 17: First version

m	n	nvdim length	dram threads	nvdim threads	dram average	dram min	dram max	nvdim average	nvdim min	nvdim max	total average	total min	total max
1000	1000000	296	10	6	0.3992	0.3424	0.4453	0.9103	0.7191	0.7623	0.9104	0.7192	0.7624
1000	1000000	246	11	5	0.4083	0.3472	0.4491	0.9008	0.7195	0.7580	0.9008	0.7195	0.7581
1000	1000000	195	12	4	0.4150	0.3683	0.4660	0.7986	0.7002	0.7149	0.7986	0.7002	0.7149
1000	1000000	146	13	3	0.4217	0.3638	0.4633	0.7781	0.6996	0.7418	0.7781	0.6996	0.7418
1000	1000000	95	14	2	0.4272	0.3806	0.4646	0.7526	0.6819	0.7276	0.7526	0.6819	0.7276
1000	1000000	48	15	1	0.4330	0.3900	0.4684	0.7405	0.6767	0.7511	0.7405	0.6767	0.7511

Table 18: Second version

6.1 First version

There are two groups of threads that works in parallel in this program. The first group of threads works on the part of the data that is stored on DRAM and the other works on the data stored on NVDIMM. One thread in each group works on data that borders with the other group. In the DRAM group that will be the last thread and in the NVDIMM group that would be the first group.

Listing 11: First version

```
1 while(k<K_length){
2     #pragma omp barrier
3     #pragma omp single
4     {
5         total_time[k] = mysecond();
6     }
7     #pragma omp barrier
8     if( thread_id < dram_threads ){
9         //for the thread bordering on nvdimmm thread.
10        if( thread_id==(dram_threads-1) ){
11            individual_time[thread_id] = mysecond();
12            for( i=slice_start; i<slice_end-1; i++){
13                for( j=1; j<nMinusOne; j++){
14                    temp = A[i-1][j-1] + A[i-1][j] + A[i-1][j+1]+
15                        A[i][j-1] + A[i][j+1]+
16                        A[i+1][j-1] + A[i+1][j] + A[i+1][j+1];
17                    B[i][j] = temp*inverseEigth;
18                }
19            }
20
21            i = slice_end-1;
22            for( j=1; j<nMinusOne; j++){
23                temp = A[i-1][j-1] + A[i-1][j] + A[i-1][j+1]+
24                    A[i][j-1] + A[i][j+1]+
25                    D_RO(C)[i*n+j] + D_RO(C)[i*n+j] +
26                    D_RO(C)[i*n+j];
27                B[i][j] = temp*inverseEigth;
28            }
29            individual_time[thread_id] = mysecond() -
30                individual_time[thread_id];
31        }else{
32            //For all the threads not bordering with nvdimmm.
33            individual_time[thread_id] = mysecond();
```

```

32     for( i=slice_start; i<slice_end; i++){
33         for( j=1; j<nMinusOne; j++){
34             temp = A[i-1][j-1] + A[i-1][j] + A[i-1][j+1]+
35                   A[i][j-1]      +      A[i][j+1]+
36                   A[i+1][j-1] + A[i+1][j] + A[i+1][j+1];
37             B[i][j] = temp*inverseEigth;
38         }
39     }
40     individual_time[thread_id] = mysecond() -
        individual_time[thread_id];
41 }
42 }else{
43     if( thread_id==dram_threads ){
44         individual_time[thread_id] = mysecond();
45         i=0;
46         for( j=1; j<nMinusOne; j++){
47             temp =
48                 A[dram_part-1][j-1]+A[dram_part-1][j]+A[dram_part-1][j+1]+
49                 D_RO(C)[i*n+(j-1)]      +
50                 D_RO(C)[i*n+(j+1)]+
51                 D_RO(C)[(i+1)*n+(j-1)] + D_RO(C)[(i+1)*n+j] +
52                 D_RO(C)[(i+1)*n+(j+1)];
53             D_RW(D)[i*n+j] = temp*inverseEigth;
54         }
55         for( i=slice_start+1; i<slice_end-1; i++){
56             for( j=1; j<nMinusOne; j++){
57                 temp = D_RO(C)[(i-1)*n+(j-1)] +
58                       D_RO(C)[(i-1)*n+j] + D_RO(C)[(i-1)*n+(j+1)]+
59                       D_RO(C)[i*n+(j-1)]      +
60                       D_RO(C)[i*n+(j+1)]+
61                       D_RO(C)[(i+1)*n+(j+1)] +
62                       D_RO(C)[(i+1)*n+j] +
63                       D_RO(C)[(i+1)*n+(j+1)];
64                 D_RW(D)[i*n+j] = temp*inverseEigth;
65             }
66         }
67         individual_time[thread_id] = mysecond() -
            individual_time[thread_id];
68     }else{
69         individual_time[thread_id] = mysecond();
70         for( i=slice_start; i<slice_end; i++){
71             for( j=1; j<nMinusOne; j++){
72                 temp = D_RO(C)[(i-1)*n+(j-1)] +

```

```

66         D_RO(C) [(i-1)*n+j] + D_RO(C) [(i-1)*n+(j+1)] +
        D_RO(C) [i*n+(j-1)] +
        D_RO(C) [i*n+(j+1)] +
67         D_RO(C) [(i+1)*n+(j-1)] +
        D_RO(C) [(i+1)*n+j] +
        D_RO(C) [(i+1)*n+(j+1)];
68     D_RW(D) [i*n+j] = temp*inverseEigth;
69     }
70 }
71     individual_time[thread_id] = mysecond() -
        individual_time[thread_id];
72 }
73 }
74 #pragma omp barrier
75 #pragma omp single
76 {
77     total_time[k] = mysecond() - total_time[k];
78     dram_time[k]=individual_time[0];
79     for(i=1;i<dram_threads;i++){
80         if(dram_time[k]<individual_time[i])
81             dram_time[k]=individual_time[i];
82     }
83     nvdimmm_time[k]=individual_time[dram_threads];
84     for(i=dram_threads+1;i<dram_threads+nvdimmm_threads;i++){
85         if(nvdimmm_time[k]<individual_time[i])
86             nvdimmm_time[k]=individual_time[i];
87     }
88     k++;
89 }
90 #pragma omp barrier
91 } //End of while

```

6.2 Second version

Listing 12: Second version

```

1 while(k<K_length){
2     #pragma omp barrier
3     #pragma omp single
4     {
5         total_time[k] = mysecond();
6     }

```

```

7  #pragma omp barrier
8  if( thread_id < dram_threads ){
9      individual_time[thread_id] = mysecond();
10     for( i=slice_start; i<slice_end; i++){
11         for( j=1; j<nMinusOne; j++){
12             temp = A[i-1][j-1] + A[i-1][j] + A[i-1][j+1]+
13                 A[i][j-1] + A[i][j+1]+
14                 A[i+1][j-1] + A[i+1][j] + A[i+1][j+1];
15             B[i][j] = temp*inverseEigth;
16         }
17     }
18     individual_time[thread_id] = mysecond() -
        individual_time[thread_id];
19 }else{
20     individual_time[thread_id] = mysecond();
21     for( i=slice_start; i<slice_end; i++){
22         for( j=1; j<nMinusOne; j++){
23             temp = D_RO(C)[(i-1)*n+(j-1)] +
24                 D_RO(C)[(i-1)*n+j] + D_RO(C)[(i-1)*n+(j+1)]+
25                 D_RO(C)[i*n+(j-1)] +
26                 D_RO(C)[i*n+(j+1)]+
27                 D_RO(C)[(i+1)*n+(j-1)] + D_RO(C)[(i+1)*n+j]
28                 + D_RO(C)[(i+1)*n+(j+1)];
29             D_RW(D)[i*n+j] = temp*inverseEigth;
30         }
31     }
32     individual_time[thread_id] = mysecond() -
        individual_time[thread_id];
33 }
34 #pragma omp barrier
35 #pragma omp single
36 {
37     total_time[k] = mysecond() - total_time[k];
38     dram_time[k]=individual_time[0];
39     for(i=1;i<dram_threads;i++){
40         if(dram_time[k]<individual_time[i])
41             dram_time[k]=individual_time[i];
42     }
43     nvdimmm_time[k]=individual_time[dram_threads];
44     for(i=dram_threads+1;i<dram_threads+nvdimmm_threads;i++){
45         if(nvdimmm_time[k]<individual_time[i])
46             nvdimmm_time[k]=individual_time[i];
47     }
48 }

```

```

45     k++;
46 }
47 #pragma omp barrier
48 }//End of while

```

		dram	dram	dram	dram	total	total	total							
m	n	threads	average	min	max	average	min	max							
2000	500000	16	0.3957	0.3727	0.4556	0.3583	0.3727	0.4556							
0.3980	0.4009	0.4019	0.4002	0.4008	0.3996	0.3996	0.3985	0.4017	0.4021	0.4001	0.4004	0.3998	0.4007	0.4016	0.4011
0.3960	0.3989	0.4044	0.3982	0.4168	0.3974	0.3977	0.3967	0.3996	0.3998	0.3981	0.3981	0.3976	0.3987	0.3995	0.3990
0.3767	0.3798	0.4116	0.4453	0.4556	0.4198	0.3775	0.3837	0.3803	0.3806	0.3845	0.3794	0.3782	0.3846	0.3860	0.3796
0.4200	0.4542	0.4555	0.4154	0.3878	0.3796	0.3740	0.3727	0.3759	0.3768	0.3806	0.3745	0.3738	0.3811	0.3755	0.3756
0.4421	0.4166	0.3869	0.3878	0.3856	0.3838	0.3844	0.3831	0.3862	0.3870	0.3849	0.3853	0.3846	0.3856	0.3861	0.4287
0.3941	0.3960	0.4019	0.3955	0.3969	0.3991	0.3952	0.3947	0.3973	0.3978	0.3961	0.3962	0.3957	0.3967	0.4188	0.4104
0.3819	0.3893	0.3863	0.3882	0.3887	0.3828	0.3828	0.3819	0.3853	0.3858	0.3839	0.3844	0.3960	0.4323	0.4407	0.4164
0.3832	0.3855	0.3878	0.3900	0.3861	0.3897	0.3841	0.3838	0.3873	0.3878	0.3901	0.4097	0.4364	0.4274	0.3977	0.3861
0.3915	0.3885	0.3906	0.3883	0.3893	0.3877	0.3876	0.3868	0.3902	0.3906	0.4187	0.4331	0.4146	0.3888	0.3900	0.3892
0.3942	0.3974	0.3991	0.3971	0.3980	0.3959	0.3964	0.3956	0.3985	0.4031	0.3973	0.4151	0.3966	0.3978	0.3985	0.3978
0.3980	0.4009	0.4020	0.4002	0.4008	0.3996	0.3996	0.3985	0.4017	0.4021	0.4001	0.4004	0.3998	0.4007	0.4016	0.4011
0.3960	0.3989	0.4044	0.3982	0.4168	0.3974	0.3977	0.3967	0.3996	0.3998	0.3981	0.3981	0.3976	0.3987	0.3995	0.3990
0.3767	0.3798	0.4116	0.4453	0.4556	0.4198	0.3775	0.3837	0.3803	0.3806	0.3845	0.3794	0.3782	0.3846	0.3860	0.3796
0.4200	0.4542	0.4555	0.4154	0.3878	0.3796	0.3740	0.3727	0.3759	0.3768	0.3806	0.3745	0.3738	0.3811	0.3755	0.3756
0.4421	0.4166	0.3869	0.3878	0.3856	0.3838	0.3844	0.3831	0.3862	0.3870	0.3849	0.3853	0.3846	0.3856	0.3861	0.4287
0.3941	0.3960	0.4019	0.3955	0.3969	0.3991	0.3952	0.3947	0.3973	0.3978	0.3961	0.3962	0.3957	0.3967	0.4188	0.4104
0.3819	0.3893	0.3863	0.3882	0.3887	0.3828	0.3828	0.3819	0.3853	0.3858	0.3839	0.3844	0.3960	0.4323	0.4407	0.4164
0.3832	0.3855	0.3878	0.3900	0.3861	0.3897	0.3841	0.3838	0.3873	0.3878	0.3901	0.4097	0.4364	0.4274	0.3977	0.3861
0.3915	0.3885	0.3906	0.3883	0.3893	0.3877	0.3876	0.3868	0.3902	0.3906	0.4187	0.4331	0.4146	0.3888	0.3900	0.3892
0.3942	0.3974	0.3991	0.3971	0.3980	0.3959	0.3964	0.3956	0.3985	0.4031	0.3973	0.4151	0.3966	0.3978	0.3985	0.3978

Table 19: First version, dram only

m:1000	n:1000000	<u>nvdimm</u>	size:48														
		<u>nvdimm</u>	dram	<u>nvdimm</u>	dram	dram	dram	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	total	total	total				
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max				
1000	1000000	296	10	6	0.4122	0.3445	0.4545	0.9137	0.7143	0.7568	0.9138	0.7143	0.7568				
0.3078	0.3090	0.2965	0.3148	0.3173	0.3173	0.3156	0.3147	0.3150	0.3357	2.4630	2.6225	2.6307	2.6286	2.6321	2.6078		
0.4292	0.3480	0.2647	0.2862	0.3140	0.4524	0.4440	0.4467	0.4473	0.4545	0.6976	0.7568	0.7558	0.7556	0.7562	0.7427		
0.3301	0.3310	0.3202	0.3402	0.3397	0.3406	0.3370	0.3380	0.3381	0.3445	0.6661	0.7234	0.7196	0.7233	0.7192	0.7076		
0.3165	0.3172	0.3275	0.4126	0.4360	0.3275	0.3235	0.3247	0.3249	0.3311	0.6576	0.7152	0.7162	0.7164	0.7204	0.7038		
0.4152	0.4485	0.4224	0.3657	0.3132	0.3192	0.3103	0.3113	0.3115	0.3174	0.6613	0.7180	0.7155	0.7159	0.7160	0.7050		
0.3587	0.3255	0.3145	0.3352	0.3348	0.3357	0.3320	0.3329	0.3331	0.3855	0.6565	0.7139	0.7114	0.7119	0.7144	0.6988		
0.3159	0.3166	0.3056	0.3262	0.3256	0.3262	0.3224	0.3545	0.4240	0.4040	0.6634	0.7230	0.7179	0.7175	0.7185	0.7098		
0.3163	0.3169	0.3062	0.3264	0.3257	0.3270	0.3953	0.4165	0.3479	0.3305	0.6561	0.7117	0.7143	0.7141	0.7124	0.7011		
0.3146	0.3153	0.3044	0.3364	0.3870	0.4390	0.3747	0.3222	0.3250	0.3286	0.6592	0.7148	0.7169	0.7169	0.7154	0.7019		
0.3120	0.3239	0.3759	0.4376	0.3846	0.3344	0.3197	0.3204	0.3209	0.3270	0.6585	0.7147	0.7129	0.7131	0.7182	0.7028		
		<u>nvdimm</u>	dram	<u>nvdimm</u>	dram	dram	dram	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	total	total	total				
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max				
1000	1000000	246	11	5	0.4045	0.3475	0.4516	0.8973	0.7128	0.7561	0.8973	0.7128	0.7561				
0.3190	0.3236	0.3247	0.3218	0.3171	0.3235	0.3191	0.3139	0.3149	0.3148	0.3351	2.3276	2.4411	2.4358	2.4330	2.4191		
0.4494	0.4509	0.4516	0.3960	0.2946	0.3014	0.2966	0.2911	0.2919	0.3382	0.4341	0.6935	0.7507	0.7508	0.7527	0.7389		
0.3420	0.3433	0.3475	0.3438	0.3393	0.3452	0.3413	0.3363	0.3373	0.3369	0.3421	0.6751	0.7314	0.7325	0.7332	0.7210		
0.3335	0.3349	0.3394	0.3355	0.3304	0.3368	0.3329	0.3269	0.3671	0.3858	0.3329	0.6550	0.7122	0.7092	0.7162	0.6999		
0.3223	0.3238	0.3286	0.3243	0.3188	0.3443	0.3794	0.4109	0.3938	0.3297	0.3306	0.6576	0.7105	0.7102	0.7128	0.6981		
0.3183	0.3198	0.3245	0.3524	0.4056	0.4203	0.3745	0.3351	0.3129	0.3123	0.3180	0.6584	0.7103	0.7116	0.7150	0.7035		
0.3249	0.3553	0.4027	0.4080	0.3487	0.3284	0.3288	0.3187	0.3199	0.3195	0.3251	0.6584	0.7166	0.7114	0.7113	0.6982		
0.4309	0.4103	0.3677	0.3209	0.3164	0.3226	0.3182	0.3130	0.3142	0.3137	0.3617	0.6597	0.7561	0.7141	0.7161	0.7026		
0.3276	0.3267	0.3273	0.3234	0.3176	0.3247	0.3202	0.3278	0.3757	0.4246	0.4096	0.6602	0.7141	0.7123	0.7086	0.6972		
0.3174	0.3190	0.3239	0.3196	0.3186	0.3579	0.3982	0.4300	0.3807	0.3300	0.3160	0.6578	0.7113	0.7149	0.7129	0.6977		
		<u>nvdimm</u>	dram	<u>nvdimm</u>	dram	dram	dram	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	total	total	total				
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max				
1000	1000000	195	12	4	0.4194	0.3610	0.4609	0.8003	0.6938	0.7376	0.8003	0.6938	0.7376				
0.3339	0.3425	0.3327	0.3383	0.3440	0.3444	0.3423	0.3390	0.3386	0.3393	0.3393	0.3510	1.5855	1.6687	1.6728	1.6619		
0.4408	0.4505	0.3848	0.3175	0.3247	0.3249	0.3223	0.3189	0.3182	0.3190	0.4399	0.4609	0.6780	0.7376	0.7374	0.7247		
0.3505	0.3593	0.3498	0.3551	0.3607	0.3610	0.3587	0.3556	0.3554	0.3559	0.3560	0.3582	0.6424	0.6976	0.6991	0.6834		
0.3319	0.3408	0.3305	0.3361	0.3426	0.3429	0.3403	0.3855	0.4219	0.4545	0.3533	0.3399	0.6414	0.7010	0.7011	0.6891		
0.3199	0.3287	0.3183	0.3239	0.3864	0.4370	0.4536	0.4050	0.3639	0.3253	0.3250	0.3280	0.6411	0.7030	0.7038	0.6879		
0.3340	0.3909	0.4350	0.4574	0.4080	0.3506	0.3294	0.3259	0.3255	0.3263	0.3262	0.3284	0.6421	0.7035	0.7026	0.6897		
0.4529	0.4074	0.3426	0.3353	0.3416	0.3419	0.3395	0.3362	0.3357	0.3364	0.3362	0.3890	0.6393	0.6938	0.6938	0.6814		
0.3506	0.3591	0.3494	0.3549	0.3607	0.3610	0.3586	0.3556	0.3552	0.3558	0.3557	0.3581	0.6421	0.6952	0.6962	0.6845		
0.3416	0.3503	0.3400	0.3462	0.3521	0.3524	0.3498	0.3468	0.3463	0.3470	0.3989	0.4004	0.6394	0.6946	0.6980	0.6814		
0.3326	0.3416	0.3311	0.3369	0.3434	0.3437	0.3413	0.3433	0.3897	0.4417	0.3980	0.3499	0.6483	0.6974	0.6950	0.6823		

Table 20: First version, more detailed 1

		<u>nvdim</u>	<u>dram</u>	<u>nvdim</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdim</u>	<u>nvdim</u>	<u>nvdim</u>	<u>total</u>	<u>total</u>	<u>total</u>		
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max		
1000	1000000	146	13	3	0.4255	0.4131	0.4619	0.7816	0.6928	0.7811	0.7816	0.6928	0.7811		
0.3501	0.3519	0.3500	0.3504	0.3500	0.3506	0.3512	0.3487	0.3498	0.3488	0.3497	0.3478	0.3600	1.3751	1.4686	1.4489
0.3158	0.3696	0.4255	0.4501	0.4500	0.4503	0.4236	0.3141	0.3149	0.3142	0.3151	0.3137	0.3251	0.7087	0.7811	0.7682
0.4619	0.4170	0.3588	0.3396	0.3412	0.3400	0.3407	0.3380	0.3390	0.3378	0.3387	0.3372	0.3793	0.6535	0.7060	0.6919
0.3443	0.3465	0.3444	0.3448	0.3447	0.3452	0.3458	0.3430	0.3466	0.3430	0.3867	0.4325	0.4108	0.6426	0.6948	0.6811
0.3460	0.3502	0.3457	0.3463	0.3461	0.3465	0.3472	0.3444	0.3792	0.4135	0.4033	0.3580	0.3486	0.6430	0.6945	0.6855
0.3462	0.3483	0.3463	0.3467	0.3463	0.3470	0.3477	0.3983	0.4137	0.3779	0.3454	0.3441	0.3492	0.6434	0.6928	0.6871
0.3438	0.3482	0.3439	0.3441	0.3507	0.4008	0.4463	0.3891	0.3435	0.3422	0.3433	0.3417	0.3468	0.6430	0.6947	0.6854
0.3471	0.3490	0.3471	0.3829	0.4342	0.3900	0.3484	0.3456	0.3467	0.3455	0.3463	0.3450	0.3501	0.6420	0.6931	0.6864
0.3474	0.3675	0.4131	0.4088	0.3547	0.3479	0.3486	0.3462	0.3471	0.3460	0.3469	0.3453	0.3505	0.6427	0.6955	0.6845
0.4163	0.4296	0.3793	0.3465	0.3470	0.3436	0.3443	0.3415	0.3427	0.3418	0.3422	0.3410	0.3575	0.6423	0.6950	0.6825
m	n	<u>nvdim</u>	<u>dram</u>	<u>nvdim</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdim</u>	<u>nvdim</u>	<u>nvdim</u>	<u>total</u>	<u>total</u>	<u>total</u>		
		length	threads	threads	average	min	max	average	min	max	average	min	max		
1000	1000000	95	14	2	0.4314	0.3890	0.4787	0.7298	0.6581	0.7028	0.7298	0.6581	0.7028		
0.3621	0.3699	0.3672	0.3706	0.3715	0.3663	0.3618	0.3581	0.3588	0.3578	0.3587	0.3593	0.3708	0.3880	1.2456	1.3069
0.4731	0.4787	0.4398	0.3529	0.3537	0.3480	0.3431	0.3395	0.3401	0.3391	0.3398	0.3405	0.3529	0.3651	0.6633	0.7028
0.3696	0.3774	0.3747	0.3778	0.3788	0.3740	0.3693	0.3659	0.3666	0.3656	0.3664	0.3668	0.3782	0.3890	0.6427	0.6620
0.3611	0.3690	0.3678	0.3696	0.3704	0.3676	0.3606	0.3573	0.3581	0.3570	0.3578	0.3584	0.4079	0.4649	0.6263	0.6631
0.3503	0.3591	0.3561	0.3596	0.3608	0.3602	0.3496	0.3458	0.3467	0.3454	0.3970	0.4437	0.4180	0.3853	0.6250	0.6626
0.3539	0.3639	0.3608	0.3644	0.3654	0.3594	0.3532	0.3618	0.3974	0.4310	0.4022	0.3613	0.3645	0.3766	0.6269	0.6606
0.3485	0.3585	0.3553	0.3592	0.3602	0.3758	0.4030	0.4197	0.4023	0.3800	0.3443	0.3446	0.3594	0.3713	0.6258	0.6581
0.3471	0.3609	0.3531	0.3872	0.4349	0.4314	0.3936	0.3632	0.3435	0.3422	0.3432	0.3438	0.3567	0.3683	0.6273	0.6616
0.3509	0.3995	0.4423	0.4271	0.3839	0.3551	0.3501	0.3497	0.3471	0.3463	0.3471	0.3505	0.3598	0.3713	0.6277	0.6601
0.4029	0.4221	0.3778	0.3667	0.3677	0.3671	0.3570	0.3535	0.3540	0.3535	0.3540	0.3545	0.3668	0.3779	0.6303	0.6601
m	n	<u>nvdim</u>	<u>dram</u>	<u>nvdim</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdim</u>	<u>nvdim</u>	<u>nvdim</u>	<u>total</u>	<u>total</u>	<u>total</u>		
		length	threads	threads	average	min	max	average	min	max	average	min	max		
1000	1000000	48	15	1	0.4496	0.4140	0.4759	0.6884	0.6252	0.6792	0.6884	0.6252	0.6792		
0.3785	0.3860	0.3854	0.3887	0.3863	0.3880	0.3793	0.3758	0.3740	0.3748	0.3755	0.3811	0.3803	0.3793	0.3827	1.1083
0.3553	0.3670	0.3637	0.3676	0.3649	0.3667	0.4085	0.4368	0.4598	0.4620	0.3785	0.3584	0.3572	0.3559	0.3527	0.6792
0.3600	0.3680	0.3695	0.4079	0.4318	0.4759	0.4100	0.3823	0.3547	0.3556	0.3564	0.3630	0.3618	0.3608	0.3582	0.6642
0.3969	0.4410	0.4579	0.4286	0.4240	0.3661	0.3564	0.3503	0.3487	0.3495	0.3500	0.3566	0.3552	0.3568	0.3522	0.6265
0.4140	0.4016	0.3820	0.3852	0.3829	0.3847	0.3755	0.3720	0.3703	0.3709	0.3717	0.3775	0.3767	0.3758	0.3731	0.6252
0.3827	0.3831	0.3813	0.3872	0.3797	0.3815	0.3716	0.3681	0.3665	0.3672	0.3679	0.3741	0.3733	0.4107	0.4442	0.6310
0.3559	0.3625	0.3619	0.3700	0.3632	0.3690	0.3540	0.3498	0.3481	0.3487	0.3907	0.4322	0.4588	0.4232	0.3849	0.6328
0.3519	0.3609	0.3644	0.3642	0.3615	0.3633	0.3519	0.3859	0.4213	0.4580	0.4261	0.3957	0.3602	0.3522	0.3487	0.6417
0.3509	0.3635	0.3582	0.3622	0.4074	0.4513	0.4617	0.4244	0.3853	0.3457	0.3464	0.3534	0.3523	0.3512	0.3484	0.6401
0.3706	0.4205	0.4451	0.4747	0.4305	0.3814	0.3527	0.3482	0.3467	0.3541	0.3548	0.3548	0.3540	0.3529	0.3500	0.6350

Table 21: First version, more detailed 2.

		<u>nvdimm</u>	<u>dram</u>	<u>nvdimm</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>total</u>	<u>total</u>	<u>total</u>
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max
1,000	1,000,000	100	10	6	0.4823	0.4213	0.5350	0.3320	0.2711	0.8619	0.5260	0.4213	0.8619
1,000	1,000,000	100	11	5	0.4628	0.3977	0.5020	0.4063	0.3391	0.9953	0.5226	0.4040	0.9953
1,000	1,000,000	100	12	4	0.4450	0.4045	0.4901	0.4728	0.4195	0.9261	0.4986	0.4218	0.9261
1,000	1,000,000	100	13	3	0.4452	0.3841	0.5099	0.5823	0.5246	1.0428	0.5823	0.5246	1.0428
1,000	1,000,000	100	14	2	0.4440	0.3814	0.4807	0.8134	0.7380	1.3800	0.8134	0.7381	1.3801
1,000	1,000,000	100	15	1	0.4367	0.3729	0.4590	1.2964	1.1864	2.2079	1.2964	1.1864	2.2080
		<u>nvdimm</u>	<u>dram</u>	<u>nvdimm</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>total</u>	<u>total</u>	<u>total</u>
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max
1,000	1,000,000	200	10	6	0.4293	0.3801	0.5221	0.6604	0.5321	1.7671	0.6604	0.5321	1.7671
1,000	1,000,000	200	11	5	0.4225	0.3513	0.4817	0.7697	0.6238	2.0151	0.7698	0.6238	2.0151
1,000	1,000,000	200	12	4	0.4146	0.3415	0.4519	0.8611	0.7571	1.7370	0.8611	0.7571	1.7370
1,000	1,000,000	200	13	3	0.3970	0.3359	0.4278	1.0750	0.9597	1.9889	1.0750	0.9597	1.9889
1,000	1,000,000	200	14	2	0.3902	0.3389	0.4095	1.5548	1.3768	2.8569	1.5548	1.3769	2.8570
1,000	1,000,000	200	15	1	0.3812	0.3273	0.3937	2.4958	2.2860	4.2250	2.4958	2.2860	4.2250
		<u>nvdimm</u>	<u>dram</u>	<u>nvdimm</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>total</u>	<u>total</u>	<u>total</u>
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max
1,000	1,000,000	300	10	6	0.4187	0.3337	0.4504	0.9607	0.7672	2.5984	0.9607	0.7672	2.5984
1,000	1,000,000	300	11	5	0.3860	0.3149	0.4189	1.1095	0.8941	2.9522	1.1095	0.8941	2.9522
1,000	1,000,000	300	12	4	0.3716	0.3147	0.4021	1.2554	1.0973	2.5936	1.2555	1.0973	2.5936
1,000	1,000,000	300	13	3	0.3553	0.2994	0.3754	1.5950	1.4262	2.9722	1.5950	1.4262	2.9722
1,000	1,000,000	300	14	2	0.3586	0.2939	0.4467	2.2757	2.0535	4.0482	2.2757	2.0536	4.0482
1,000	1,000,000	300	15	1	0.3440	0.2899	0.4570	3.6785	3.3726	5.9822	3.6785	3.3726	5.9823
		<u>nvdimm</u>	<u>dram</u>	<u>nvdimm</u>	<u>dram</u>	<u>dram</u>	<u>dram</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>total</u>	<u>total</u>	<u>total</u>
m	n	length	threads	threads	average	min	max	average	min	max	average	min	max
1,000	1,000,000	400	10	6	0.3859	0.3285	0.4935	1.2587	0.9947	1.0591	1.2587	0.9947	1.0591
1,000	1,000,000	400	11	5	0.3522	0.3021	0.4969	1.4667	1.1868	1.2097	1.4667	1.1868	1.2097
1,000	1,000,000	400	12	4	0.3175	0.2716	0.4268	1.6739	1.4492	1.5125	1.6739	1.4492	1.5126
1,000	1,000,000	400	13	3	0.3242	0.2898	0.4059	2.1027	1.8852	1.9380	2.1028	1.8852	1.9380
1,000	1,000,000	400	14	2	0.2966	0.2831	0.4023	3.0072	2.7257	2.8082	3.0072	2.7257	2.8083
1,000	1,000,000	400	15	1	0.2930	0.2775	0.3902	4.9659	4.5704	5.0607	4.9659	4.5704	5.0607

Table 22: First version.

		nvdi	dram	nvdi	dram	dram	dram	nvdim	nvdim	nvdim	total	total	total
m	n	leng	threa	threa	average	min	max	average	min	max	average	min	max
1,000	1,000,000	100	10	6	0.5031	0.4696	0.5522	0.3221	0.2608	0.2654	0.5421	0.4696	0.5522
1,000	1,000,000	100	11	5	0.4612	0.4093	0.4975	0.3968	0.3246	0.3328	0.5246	0.4093	0.4975
1,000	1,000,000	100	12	4	0.4513	0.4098	0.4884	0.4552	0.4048	0.4098	0.5014	0.4098	0.4884
1,000	1,000,000	100	13	3	0.4422	0.3912	0.4799	0.5642	0.5084	0.5379	0.5642	0.5084	0.5379
1,000	1,000,000	100	14	2	0.4334	0.3799	0.4681	0.8098	0.7294	0.8177	0.8098	0.7295	0.8177
1,000	1,000,000	100	15	1	0.4367	0.4298	0.4616	1.4619	1.3356	1.4728	1.4619	1.3356	1.4728
		nvdi	dram	nvdi	dram	dram	dram	nvdim	nvdim	nvdim	total	total	total
m	n	leng	threa	threa	average	min	max	average	min	max	average	min	max
1,000	1,000,000	200	10	6	0.4462	0.4150	0.5193	0.6424	0.5142	0.5365	0.6424	0.5142	0.5365
1,000	1,000,000	200	11	5	0.4145	0.3720	0.4703	0.7537	0.6025	0.6358	0.7537	0.6025	0.6358
1,000	1,000,000	200	12	4	0.4130	0.3786	0.4469	0.8285	0.7306	0.7617	0.8285	0.7306	0.7617
1,000	1,000,000	200	13	3	0.4032	0.3466	0.4310	1.0511	0.9367	1.0338	1.0511	0.9367	1.0338
1,000	1,000,000	200	14	2	0.3898	0.3712	0.4123	1.4831	1.3442	1.3693	1.4831	1.3442	1.3694
1,000	1,000,000	200	15	1	0.3867	0.3754	0.4716	2.8308	2.5865	2.6867	2.8309	2.5866	2.6867
		nvdi	dram	nvdi	dram	dram	dram	nvdim	nvdim	nvdim	total	total	total
m	n	leng	threa	threa	average	min	max	average	min	max	average	min	max
1,000	1,000,000	300	10	6	0.3999	0.3447	0.4524	0.9276	0.7416	0.7500	0.9276	0.7416	0.7500
1,000	1,000,000	300	11	5	0.3932	0.3244	0.4249	1.0898	0.8660	0.9322	1.0898	0.8660	0.9322
1,000	1,000,000	300	12	4	0.3749	0.3143	0.4042	1.2065	1.0654	1.0783	1.2065	1.0654	1.0783
1,000	1,000,000	300	13	3	0.3481	0.3203	0.3696	1.5408	1.3923	1.4199	1.5408	1.3923	1.4199
1,000	1,000,000	300	14	2	0.3501	0.3459	0.3733	2.1808	2.0092	2.0432	2.1808	2.0092	2.0432
1,000	1,000,000	300	15	1	0.3404	0.3264	0.4569	4.1787	3.8467	3.9795	4.1788	3.8468	3.9795
		nvdi	dram	nvdi	dram	dram	dram	nvdim	nvdim	nvdim	total	total	total
m	n	leng	threa	threa	average	min	max	average	min	max	average	min	max
1,000	1,000,000	400	10	6	0.3710	0.2957	0.5086	1.2246	0.9551	1.0135	1.2246	0.9551	1.0135
1,000	1,000,000	400	11	5	0.3437	0.2816	0.4704	1.4251	1.1364	1.2231	1.4251	1.1364	1.2231
1,000	1,000,000	400	12	4	0.3336	0.2733	0.4385	1.5954	1.3988	1.4362	1.5954	1.3988	1.4363
1,000	1,000,000	400	13	3	0.3182	0.2837	0.4174	2.0210	1.8115	1.8356	2.0210	1.8115	1.8356
1,000	1,000,000	400	14	2	0.3018	0.2896	0.4058	2.9019	2.6311	2.7827	2.9019	2.6312	2.7827
1,000	1,000,000	400	15	1	0.2974	0.2847	0.4018	5.4534	5.0678	5.1469	5.4535	5.0679	5.1469

Table 23: Second version.

m	n	<u>nvdim</u> length	dram threads	<u>nvdim</u> threads	dram average	dram min	dram max	<u>nvdim</u> average	<u>nvdim</u> min	<u>nvdim</u> max
1,000	1,000,000	200	1	15	2.4969	2.3980	2.6683	0.2652	0.1845	0.5940
1,000	1,000,000	200	2	14	1.4481	1.2292	1.7017	0.2503	0.1988	0.6042
1,000	1,000,000	200	3	13	1.0203	0.8528	1.2142	0.2675	0.2131	0.6043
1,000	1,000,000	200	4	12	1.0637	0.6941	1.2707	0.4761	0.2276	0.6246
1,000	1,000,000	200	5	11	0.7014	0.5716	0.8123	0.3013	0.2567	0.6231
1,000	1,000,000	200	6	10	0.5767	0.5035	0.6715	0.3755	0.2848	0.8516
1,000	1,000,000	200	7	9	0.5649	0.4562	0.6394	0.3539	0.3151	0.6364
1,000	1,000,000	200	8	8	0.4774	0.4289	0.5666	0.4176	0.3599	0.6759
1,000	1,000,000	200	9	7	0.4865	0.4106	0.5245	0.4371	0.4061	0.7098
1,000	1,000,000	200	10	6	0.4713	0.3851	0.5153	0.5126	0.4728	0.8159
1,000	1,000,000	200	11	5	0.4563	0.3681	0.4958	0.6027	0.5551	0.9569
1,000	1,000,000	200	12	4	0.3926	0.3576	0.4642	0.7724	0.6760	1.1671
1,000	1,000,000	200	13	3	0.3822	0.3459	0.4368	0.9688	0.8599	1.5051
1,000	1,000,000	200	14	2	0.3931	0.3430	0.4605	1.5638	1.2419	2.4540
1,000	1,000,000	200	15	1	0.3820	0.3332	0.4844	2.7315	2.3613	4.0120

Table 24: First version. OLD

m	n	<u>nvdim</u> length	dram threads	<u>nvdim</u> threads	dram average	dram min	dram max	<u>nvdim</u> average	<u>nvdim</u> min	<u>nvdim</u> max
1000	1000000	400	1	15	1.9421	1.7041	2.5587	0.5400	0.3737	1.1992
1000	1000000	400	2	14	1.6456	1.3045	2.1652	0.9189	0.7172	1.4330
1000	1000000	400	3	13	0.7794	0.6837	0.9609	0.5340	0.4342	1.1997
1000	1000000	400	4	12	0.7031	0.5605	0.9786	0.7232	0.4821	1.2280
1000	1000000	400	5	11	0.5446	0.4239	0.6813	0.6175	0.5254	1.2250
1000	1000000	400	6	10	0.4830	0.3795	0.6112	0.6968	0.5731	1.2296
1000	1000000	400	7	9	0.4217	0.3497	0.5369	0.7117	0.6213	1.2246
1000	1000000	400	8	8	0.4217	0.3202	0.6186	0.9400	0.6843	1.3209
1000	1000000	400	9	7	0.3489	0.3263	0.4202	0.8846	0.7568	1.4615
1000	1000000	400	10	6	0.3534	0.2994	0.4695	1.2051	0.9452	1.6556
1000	1000000	400	11	5	0.3400	0.3112	0.4017	1.1752	1.0287	1.8976
1000	1000000	400	12	4	0.3907	0.2700	0.5614	2.4105	1.5335	2.8562
1000	1000000	400	13	3	0.2984	0.2701	0.3220	1.8764	1.6555	2.8262
1000	1000000	400	14	2	0.2894	0.2565	0.3947	2.7862	2.4556	3.9577
1000	1000000	400	15	1	0.3067	0.2609	0.3710	5.4896	4.7125	7.1350

Table 25: First version. OLD

		<u>nvdimm</u>	dram	<u>nvdimm</u>	dram	dram	dram	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>
m	n	length	threads	threads	average	min	max	average	min	max
1,000	1,000,000	200	1	15	4.3506	3.0851	5.4226	12.5578	12.0879	13.9106
1,000	1,000,000	200	2	14	2.4177	1.6629	3.3427	12.5061	11.9543	13.8187
1,000	1,000,000	200	3	13	1.6175	1.1164	2.1824	12.0978	11.6232	13.2806
1,000	1,000,000	200	4	12	1.2294	0.8984	1.6496	11.8080	11.4248	12.8516
1,000	1,000,000	200	5	11	0.9603	0.5595	1.3311	11.2715	10.9282	12.2981
1,000	1,000,000	200	6	10	0.8311	0.4913	1.1458	11.2827	11.0103	11.8192
1,000	1,000,000	200	7	9	0.7681	0.4157	1.2062	11.0349	10.3212	12.9057
1,000	1,000,000	200	8	8	0.6855	0.4458	0.9048	10.7173	10.1667	12.3145
1,000	1,000,000	200	9	7	0.6231	0.3692	1.0769	10.5659	9.9144	12.4387
1,000	1,000,000	200	10	6	0.5536	0.3801	0.7129	10.2459	9.4983	12.5660
1,000	1,000,000	200	11	5	0.4918	0.3547	0.6574	9.5222	9.1636	10.7587
1,000	1,000,000	200	12	4	0.4674	0.3504	0.6201	8.6057	8.3691	9.3879
1,000	1,000,000	200	13	3	0.4439	0.3312	0.5859	7.9526	7.5966	8.5104
1,000	1,000,000	200	14	2	0.4377	0.3325	0.6917	5.6382	5.1366	6.2976
1,000	1,000,000	200	15	1	0.3979	0.3336	0.5328	2.7644	2.4067	3.7062

Table 26: Second version. OLD

		<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>	<u>nvdimm</u>
m	n	threads	average	min	max
200	1,000,000	1	2.9744	2.5180	3.9372
200	1,000,000	2	4.9422	4.6678	5.4835
200	1,000,000	3	5.0508	4.8935	5.6034
200	1,000,000	4	5.5549	5.3542	6.7077
200	1,000,000	5	5.2473	4.9133	6.5496
200	1,000,000	6	5.3491	5.0966	6.7279
200	1,000,000	7	5.6929	5.4649	6.3630
200	1,000,000	8	5.8173	5.3828	6.8326
200	1,000,000	9	5.3901	5.2969	5.9944
200	1,000,000	10	5.5692	5.2874	6.2419
200	1,000,000	11	5.4967	5.2975	6.0858
200	1,000,000	12	5.8591	5.6733	7.0265
200	1,000,000	13	6.0996	5.8470	7.0950
200	1,000,000	14	5.7147	5.5264	6.6942
200	1,000,000	15	5.7758	5.5657	6.6912
200	1,000,000	16	5.8444	5.6211	6.4624

Table 27: NVDIMM only of second version. OLD

				predicted	
m	n	threads	time	time	bandwidth
1,000	1,000,000	1	1.7719	0.4233	18,898.7
1,000	1,000,000	2	1.0333	0.2800	28,571.4
1,000	1,000,000	3	1.0121	0.2744	29,158.2
1,000	1,000,000	4	0.9899	0.2736	29,235.5
1,000	1,000,000	5	0.6998	0.2011	39,789.2
1,000	1,000,000	6	0.6439	0.1792	44,646.4
1,000	1,000,000	7	0.5595	0.1565	51,121.6
1,000	1,000,000	8	0.5318	0.1486	53,839.6
1,000	1,000,000	9	0.4858	0.1357	58,964.7
1,000	1,000,000	10	0.4578	0.1263	63,321.0
1,000	1,000,000	11	0.4300	0.1198	66,792.3
1,000	1,000,000	12	0.3985	0.1139	70,227.7
1,000	1,000,000	13	0.3873	0.1083	73,838.5
1,000	1,000,000	14	0.3694	0.1041	76,823.2
1,000	1,000,000	15	0.3594	0.1004	79,680.9
1000	1000000	16	0.3394	0.0972	82,326.0
Formula for prediction					
$(m*n*8*0.000001)/\text{bandwidth}$					

Table 28: DRAM only on n50. OLD

References

- [1] John D. McCalpin. *STREAM source code*. URL: <https://www.cs.virginia.edu/stream/FTP/Code/stream.c> (visited on 12/20/2020).