1 Simulation

1.1 Introduction

The tests that was done in chapter three and four the DRAM and NVDIMM have been doing the same type of work. This chapter will be about DRAM and NVDIMM doing different types of work. There will be a group of threads calculating data on DRAM. While this is going on there will be another group of threads copying the latest set of calculated data to NVDIMM and analyzing the data on NVDIMM. The purpose of this is to find out if NVDIMM can work on other tasks and understand how NVDIMM behave when working other threads are working on DRAM. The chapter will start by testing the simulation using only DRAM in order to see what kind of performance is possible. Afterward another version of the program will be tested where both DRAM and NVDIMM is being used.

1.2 DRAM only

1.2.1 Description of code

The code at start by taking a time measurement before the while-loop in line 4 in listings 1 by on of the threads. This time measurement will be ended after the while-loop at line 41 by one thread. Since there is a barrier at line 38 the time measurement will only end when all the threads have left the while-loop.

The While-loop at line 6 will be repeated for 5000 times and there is a barrier in line 7 to ensure that all threads are finished with the previous iteration before starting a new iteration. One thread will start by entering a single in line 8 and do all the work only one thread can do such as increase n by one, zero out diff and average and a part of a formula that will be used in data generation. The description of data generation will be done in section 1.2.2.

After the data generation all the threads will wait the barrier at line 19 before one thread enter the single at line 20. The thread will swap the arrays x and xk_1 before starting a new time measurement at line 26. The threads will move on to analyzing the data generated and this will be described at section 1.2.3.

The threads will wait for all the threads to complete analyzing the data at line 31 before on thread enter a new single at line 32 where it will first start by calculating an average and after that it will end the

time measurement that was started at line 26. This is the time measurement of the analyzing part of the code, the time it takes to analyze the data in all iteration will be add to one variable called analyze_time. This is the end of the while-loop and the code will either return to line 7 or exit the while-loop

The time measurement of the data generation will be found by subtracting analyze_time from the time measurement that was taken before and after the code.

Listing 1: DRAM only version of the simulation code.

```
double double_temp;
#pragma omp single
{
   data_generation_time, = mysecond();
}
while( n<5000 ){
   #pragma omp barrier
   #pragma omp single
   {
      n++;
      diff=0.0;
      average = 0;
      //completes the first part of the formula.
      Wk_1_product = (omd + (d*Wk_1))*iN;
   }

   Data generation, see section 1.2.2

   #pragma omp barrier
   #pragma omp single
   {
      temp_x = xk_1;
      xk_1 = x;
      x = temp_x;
      //starting time measurement of calculation.
      temp_calc=mysecond();
   }

   Analyzing the data, see section 1.2.3

   #pragma omp barrier
   #pragma omp single
   {
```

```
34      average *= iN;
35      analyze_time+=mysecond()-temp_calc;
36    }
37 }
38 #pragma omp barrier
39 #pragma omp single
40 {
41    data_generation_time, = mysecond() -
          data_generation_time;
42 }
```

The code shown in listing 2 is the function used to take the time measurement. This function have been copied from the Stream benchmark code.

Listing 2: Function that that is used for measure time

```
1 double mysecond(){
2    struct timeval tp;
3    struct timezone tzp;
4    int i;
5    i = gettimeofday(&tp,&tzp);
6    return ( (double) tp.tv_sec + (double) tp.tv_usec *
          1.e-6 );
7 }
```

1.2.2 Data generation

All the threads will enter the for-loop at line 2 and they will start by zero out the double_temp variable. They will then enter a new for-loop at line 5. The array the for-loop will pass through is a compressed row storage(CRS), this means that a 2D-array have been converted to an 1D-array where all the elements with zero have been removed. CRS_row_ptr shows the where a row in the 2D-array begin and end in the CRS, that is why the CRS_row_ptr is used in the header of the for-loop. At line 6 the for-loop will multiply together all elements in a row together with the result from a previous iteration of data generation. Everything will be added together in the variable double_temp. The variable double_temp will then be multiplied with a constant d at line 9 and added together with a constant Wk_1_product at line 11 before being stored in array x at line 12 which is where data from the current data generation is being stored. The code will then check if the dif-

ference between current data generation and previous data generation is higher than the diff variable. If it is then the diff variable will be changed to the current difference.

Listing 3: Generation of data.

```
1   #pragma omp for reduction(max:diff)
2   for( i=0; i<nodes; i++){
3       //This is A*x^k-1
4       double_temp = 0;
5       for( j=CRS_row_ptr[i]; j<CRS_row_ptr[i+1]; j++){
6           double_temp += CRS_values[j] * xk_1[CRS_col_idx[j]];
7       }
8       //d*Ax^k-1
9       double_temp *= d;
10      //Adding the first part and second part together.
11      double_temp += Wk_1_product;
12      x[i] = double_temp;
13
14      //Comuting the difference between x^k and x^k-1
15      //and adds the biggest diff to diffX[thread_id]
16      if( x[i]-xk_1[i] > diff ){
17          diff = x[i] - xk_1[i];
18      }
19  }
```

1.2.3 Analyze

The analyze part of the code run through the nodes array five times and add all the elements to the variable called average. The average variable is divided by the number of elements in array, this is shown in the code in listing 4

Listing 4: Analyzing the data.

```
1   //Analyse part
2   #pragma omp for reduction(+ : average)
3   for(i=0;i<nodes;i++){
4      average += xk_1[i];
5   }
6   #pragma omp for reduction(+ : average)
7   for(i=0;i<nodes;i++){
8      average += xk_1[i];
9   }
10  #pragma omp for reduction(+ : average)
11  for(i=0;i<nodes;i++){
12     average += xk_1[i];
13  }
14  #pragma omp for reduction(+ : average)
15  for(i=0;i<nodes;i++){
16     average += xk_1[i];
17  }
18  #pragma omp for reduction(+ : average)
19  for(i=0;i<nodes;i++){
20     average += xk_1[i];
21  }
```

1.2.4 Prediction

By calculating how much data is tranferred back and forth between memory and CPU it is possible to calculate how much time it will take to complete the program. The data used in this program is a text-file that contain 16,000,000 nodes and 127,952,004 edges that goes between the nodes. The arrays will have the length of either nodes or edges.

During one iteration the data generation will load data from two array with the length of edges and one is of type double and the of the type int. The array names are CRS_row_ptr and CRS_values. The way the data traffic is calculated is shown in figure 1.

5

$$\frac{Nodes * 2.5 * 5000 * 8}{1000000}$$

Figure 1: Calculation of data traffic generated by data generation.

$$\frac{edges * 1.5 * 5000 * 8}{1000000}$$

Figure 2: Calculation of data traffic generated by analyzing the data.

Arrays of the length of nodes will be load or store two times of type double and one time from type int. The array names are x, xk_1 and CRS_col_idx. The data traffic generated by analyzing data is calculated is shown in 2.

The DRAM speed comes from the DRAM benchmark in chapter 3.

Threads	DRAM speed	predicted data genration time	predicted analyze time
1	11673.50	794.72	274.13
2	22995.10	403.44	139.16
3	33554.90	276.48	95.37
4	42917.30	216.16	74.56
5	50260.90	184.58	63.67
6	53612.50	173.04	59.69
7	56100.80	165.37	57.04
8	58554.60	158.44	54.65
9	60491.70	153.36	52.90
10	62242.20	149.05	51.41
11	64257.10	144.38	49.80
12	64890.30	142.97	49.31
13	65648.80	141.31	48.74
14	65606.50	141.41	48.78
15	65665.50	141.28	48.73
16	65509.80	141.61	48.85

Table 1: Time prediction, Data generation and analyze of data

1.2.5 Result

Threads	Total time	Data generation	analyze time
1	1545.98	949.60	596.38
2	671.34	421.51	249.82
3	459.76	289.57	170.18
4	350.04	220.99	129.05
5	284.09	179.43	104.66
6	240.27	151.90	88.37
7	209.05	132.54	76.51
8	188.60	120.08	68.52
9	177.25	113.20	64.05
10	168.02	109.87	58.15
11	159.60	106.06	53.54
12	152.81	103.37	49.44
13	150.41	101.86	48.55
14	146.30	100.99	45.31
15	143.02	100.55	42.47
16	140.75	100.74	40.01

1.3 NVM simulation

1.3.1 Locks

The program are divided into two parts, the calculation of data and the analyzing of the data that have been generated. The two parts synchronize by using two locks called lock_a and lock_b, lock_a will start in unlocked state and lock_b in locked state. When the two parts starts the analyzing part is put on hold by lock_b until the calculation part has generated the first set of data. Then the calculation part will lock lock_a, swap the pointers x and xk_1 and then unlock lock_b. The calculation part will then start the calculation of the next set of data, but wont swap pointers until analyzing part has transferred the content in xk_1 to NVDIMM and unlocked lock_a.

When the calculation part unlocks lock_b the analyzing will start transferring data from xk_1 to NVDIMM and unlock lock_a when it's done with the transfer. The analyzing part will then start analyzing the data on NVDIMM. When its done it will encounter lock_b and will wait there until calculation has a new set of data ready and has swapped the pointers.

Before and after the set lock in calculation and analyze part there is a time measurement that measure how long the threads have waited for the lock to be unlocked by the other part. All the individual times the threads have waited in calculation or analyze gets added to a variable called iteration_idle_time or transfer_idle_time that will be the total time the threads have waited.
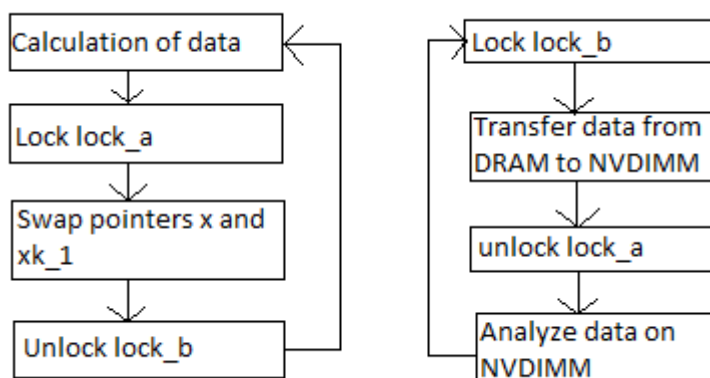

Figure 3: A simplified version of how lock works in the program.

1.3.2 Calculation

The total time for calculation are measured before and after the while-loop. One thread will start the time measurement at line 3 and the measurement will be ended by one thread at line 33. The calculation starts with a while-loop at line 5. It will be repeated 5000 times. There is a barrier at line 6 that will synchronize all the threads. One thread will enter a single bracket a line 7 and increase the variable n by one, zero out the variables diff and Wk_1 and calculate one part of the calculation that only need to be done by one thread at line 11.

The calculation of data will happens at line 17. All threads will enter the for-loop at line 18, this for-loop will pass through all the nodes in the array. A temporary variable is created at line 20, this is to ensure that data is only stored at the end of the for-loop. Each thread will then enter a second for-loop at line 21. The arrays that begins with CRS is part of a compressed sparse matrix, that means that a 2D-array has been turned into a 1D-array in order to perform faster. The for-loop will add together edges that is connected to node i, they will be stored in the double_temp variable. Once out of the array the double_temp will be multiplied with a constant d at line 25 and added together with the Wk_product at line 27, this is a variable that was calculated by one thread in line 13. The double_temp value will be written to array at line 28.

One thead will enter the single bracket at line 36 and will start a new time measurement at line 38. This time measurement will measure how much time the calculation threads must wait for the analyze threads to finish. Once the analyze threads have unlocked lock_a the thread will pass through the lock and lock it again at line 39. It will also end the time measurement at line 40 and add the time waited to the iteration_idle_time variable. This variable represent the total amount of time the calculation threads have been idle. At line 42-44 the arrays x and xk_1 will be swapped. At line 46 the thread will unlock the lock_b, this will allow the analyze threads to do their jobs. The threads will then return to the beginning of the while-loop at line 5 and repeat the process if n is lower than 5000.

Listing 5: Calculation

```
1  #pragma omp single
2  {
3      iteration_time = mysecond();
4  }
```

```
5   while( n<5000 ){
6      #pragma omp barrier
7      #pragma omp single
8      {
9         n++;
10        diff=0.0;
11        Wk_1=0;
12        //completes the first part of the formula.
13        Wk_1_product = (omd + (d*Wk_1))*iN;
14     }
15
16     /* Calculation of Data */
17     #pragma omp for reduction(max:diff)
18     for( i=0; i<nodes; i++){
19        //This is A*x^k-1
20        double_temp = 0;
21        for( j=CRS_row_ptr[i]; j<CRS_row_ptr[i+1]; j++){
22           double_temp += CRS_values[j] * xk_1[CRS_col_idx[j]];
23        }
24        //d*Ax^k-1
25        double_temp *= d;
26        //Adding the first part and second part together.
27        double_temp += Wk_1_product;
28        x[i] = double_temp;
29
30        //Comuting the difference between x^k and x^k-1
31        //and adds the biggest diff to diffX[thread_id]
32        if( x[i]-xk_1[i] > diff )
33           diff = x[i] - xk_1[i];
34     }
35
36     #pragma omp single
37     {
38        temp_time = mysecond();
39        omp_set_lock(&lock_a);
40        iteration_idle_time += mysecond() - temp_time;
41
42        temp_x = xk_1;
43        xk_1 = x;
44        x = temp_x;
45
46        omp_unset_lock(&lock_b);
47     }
```

```
48  }//end of while-loop
49  #pragma omp single
50  {
51    iteration_time = mysecond() - iteration_time;
52  }
```

1.3.3 Analyze

The time will be measured before and after the while-loop that start at line 1. The if test for this while-loop has been moved to the end of the while-loop, that is why the it test at line 1 is 1==1 and will allways be true. One thread will enter the single bracket at line 2 and start a time measurement that will measure the idle time of the analyze threads. The threads will wait until the calculation threads unlocks the lock_b. When its unlocked the thread will pass through line 5 and lock lock_b again. It will also add the time waited to the variable transfer_idle_time. The thread will then start a new time measurement at line 7, this measurement will measure the time it takes to transfer data from DRAM to NVDIMM. the variable average will be turned to zero at line 8. At line 11-14 all the threads will work together to transfer the xk_1 array from DRAM to NVDIMM. Once done one thread will enter a new single at line 15. The thread will then end the time measurement at line 17 and add the time taken to transfer the data to the variable DRAM_to_NVM_time. It will then unlock lock_a at line 18, this will allow the calculation threads to swap x and xk_1. The thread will also start a new time measurement at line 19, this measurement will measure how long it takes for the threads to analyze the data. From line 22-25 is where the data gets analyzed, in this case its just an average of all the data. The threads will synchronize at a barrier at line 26. One thread will enter the single thread at line 27 and will divide the sum of all the nodes with the number of nodes. It will then end the time measurement at line 30 and add it to the total time it takes to analyze the data. All the threads will then go through an if test at line 33. This test will always be true until the calculation threads changes the iteration_ongoing variable from 1 to 0.This will only be done when calculation threads are done with calculation and left their while-loop that was explained above.

Listing 6: Analyze

```
1  while(1==1){
```

```
2    #pragma omp single
3    {
4      temp_time = mysecond();
5      omp_set_lock(&lock_b);
6      transfer_idle_time += mysecond() - temp_time;
7      temp_time = mysecond();
8      average=0.0;
9    }
10   /* Transfer of array from DRAM to NVDIMM */
11   #pragma omp for
12   for(i=0; i<nodes; i++){
13     D_RW(nvm_values)[i]=xk_1[i];
14   }
15   #pragma omp single
16   {
17     DRAM_to_NVM_time += mysecond() - temp_time;
18     omp_unset_lock(&lock_a);
19     temp_time = mysecond();
20   }
21   /* Analyzations of data */
22   #pragma omp for reduction(+ : average)
23   for(i=0;i<nodes;i++){
24     average += D_RO(nvm_values)[i];
25   }
26   #pragma omp barrier
27   #pragma omp single
28   {
29     average /= nodes;
30     Analyse_time += mysecond() - temp_time;
31   }
32   //if sentence for exiting while-loop.
33   if(iteration_ongoing==0){
34     break;
35   }
36 }
```

1.3.4 Prediction

1.3.5 Result

Total threads	DataGen Threads	Analyze threads	DataGen time	DataGen idle time	Total Time
16	15	1	102.35	849.87	952.37
16	14	2	105.89	402.37	508.35
16	13	3	112.69	274.05	386.80
16	12	4	117.66	181.92	299.63
16	11	5	123.18	126.36	249.58
16	10	6	130.29	90.92	221.25
16	9	7	136.33	51.77	188.13
16	8	8	147.03	18.63	165.69
16	7	9	169.39	3.78	173.20
16	6	10	182.96	0.02	183.01
16	5	11	206.64	0.00	206.67
16	4	12	246.10	0.31	246.46
16	3	13	303.90	0.00	303.93
16	2	14	430.30	0.00	430.34
16	1	15	804.03	0.00	804.08

Table 2: Result of generating data.

Total threads	DataGen Threads	Analyze threads	Transfer idle time	DRAM-NVM time	Analyze time	Total Time
16	15	1	0.03	166.02	786.30	952.37
16	14	2	0.04	84.74	423.54	508.35
16	13	3	0.03	70.52	316.22	386.80
16	12	4	0.01	51.48	248.09	299.63
16	11	5	0.03	43.15	206.36	249.58
16	10	6	0.03	44.56	176.63	221.25
16	9	7	0.10	34.79	153.22	188.13
16	8	8	0.03	31.22	134.42	165.69
16	7	9	16.01	36.38	120.80	173.20
16	6	10	51.34	28.85	102.81	183.01
16	5	11	90.16	26.39	90.11	206.67
16	4	12	127.23	28.67	90.41	246.46
16	3	13	205.38	22.83	75.71	303.93
16	2	14	338.74	21.71	69.86	430.34
16	1	15	718.38	20.60	65.07	804.08

Table 3: Result of transfer data to NVDIMM and analyzing the data.